



# A Two-layer Partitioning for Non-point Spatial Data

Dimitrios Tsitsigkos<sup>1,3</sup> Konstantinos Lampropoulos<sup>3</sup> Panagiotis Bouros<sup>2</sup>  
Nikos Mamoulis<sup>3</sup> Manolis Terrovitis<sup>1</sup>

<sup>1</sup>Athena RC, Greece

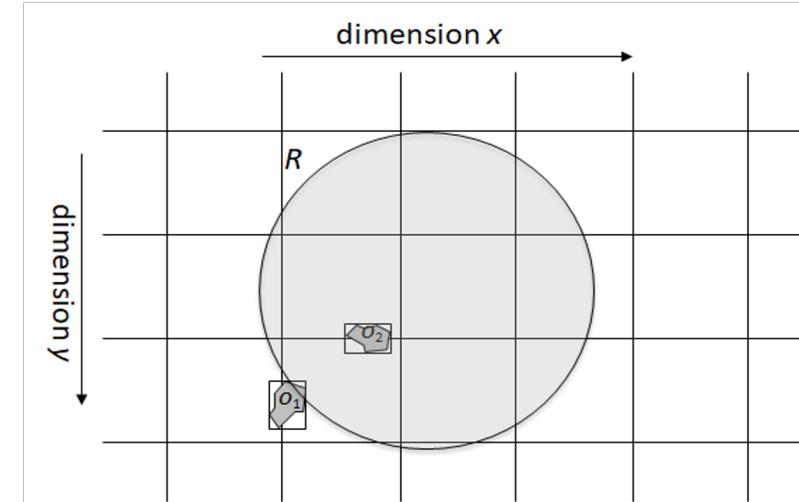
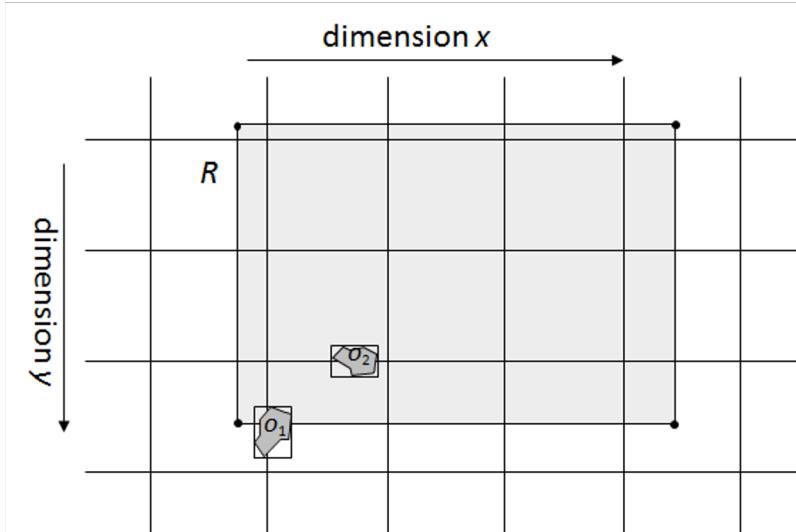
<sup>2</sup>Johannes Gutenberg University Mainz, Germany

<sup>3</sup>University Of Ioannina, Greece



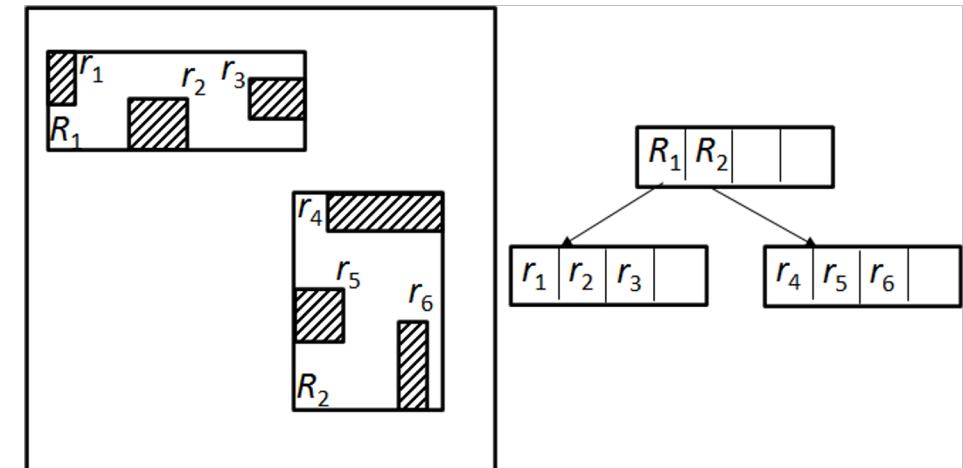
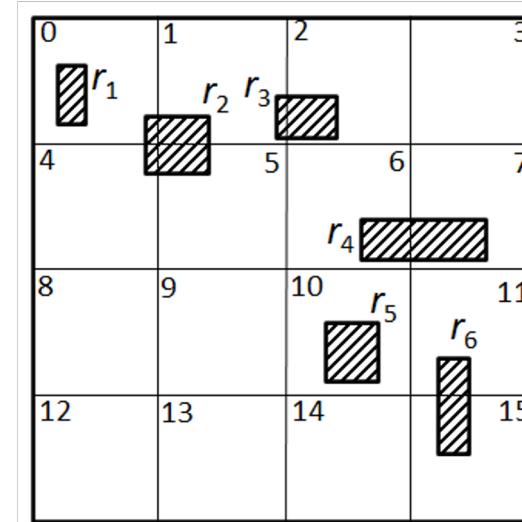
# Range Query

- Range query is a **fundamental** operation in managing spatial data:
  - **Geographic Information Systems** (e.g., management of huge meshes)
  - **Neuroscience** (e.g., building and indexing a spatial model of brain)
  - **Location-based analytics** (e.g., managing spatial influence region of mobile users in order to facilitate effective POI recomm)
- Retrieve all spatial objects which **intersect** with the **area** of  $R$



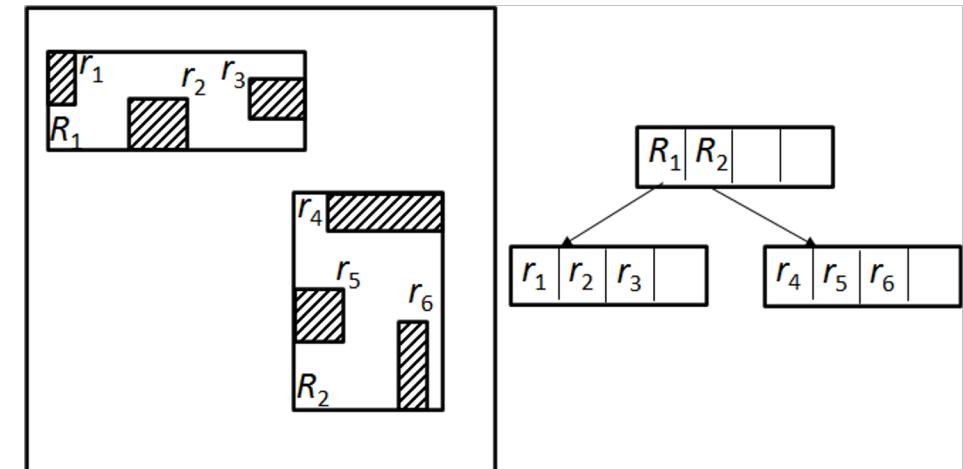
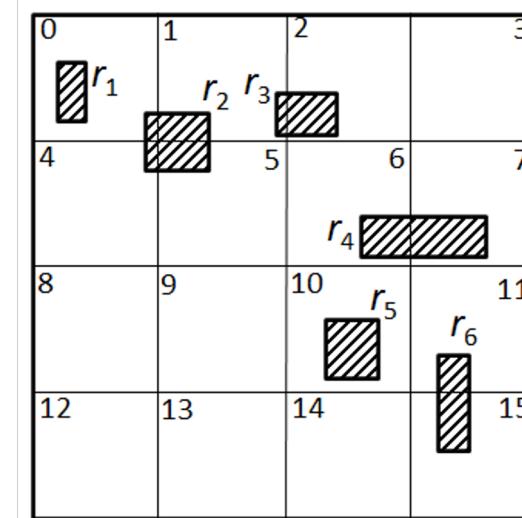
# Motivation (1/2)

- Focus on **non-point** data
- **Space**-Oriented Partitioning:
  - Divide the space into spatially **disjoint** partitions
  - Data **replication**
  - Grid, quad-tree, etc.
- **Data**-Oriented Partitioning:
  - **No** data **replication**
  - Balanced structure
  - R-tree family, etc.



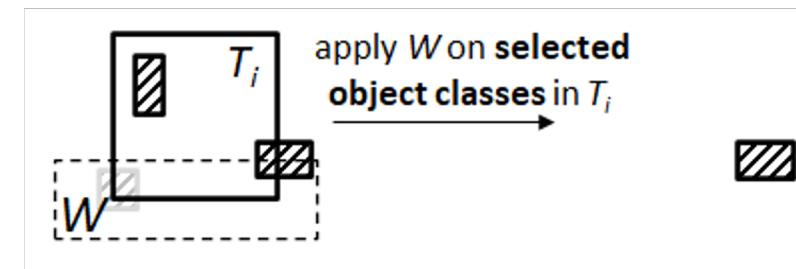
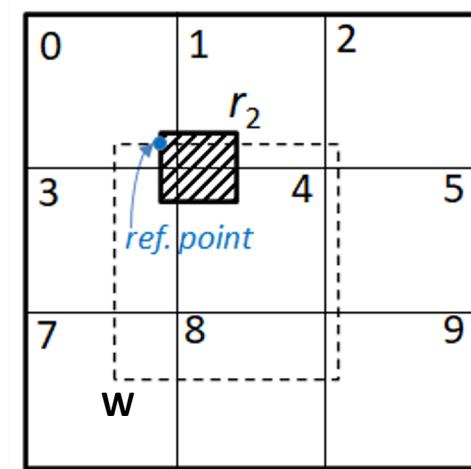
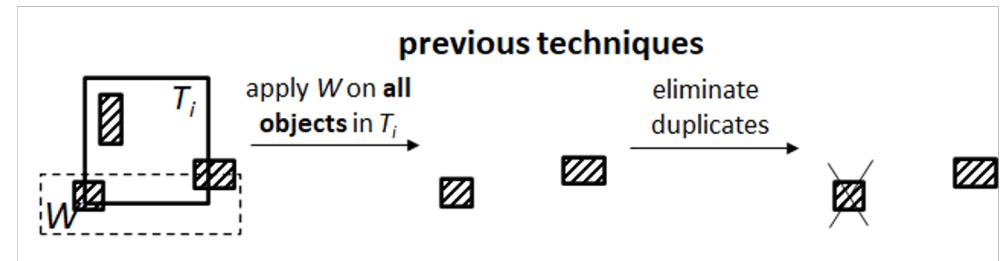
# Motivation (1/2)

- Focus on **non-point** data
- **Space**-Oriented Partitioning:
  - Divide the space into spatially **disjoint** partitions
  - Data **replication**
  - Grid, quad-tree, etc.
- **Data**-Oriented Partitioning:
  - **No** data **replication**
  - Balanced structure
  - R-tree family, etc.



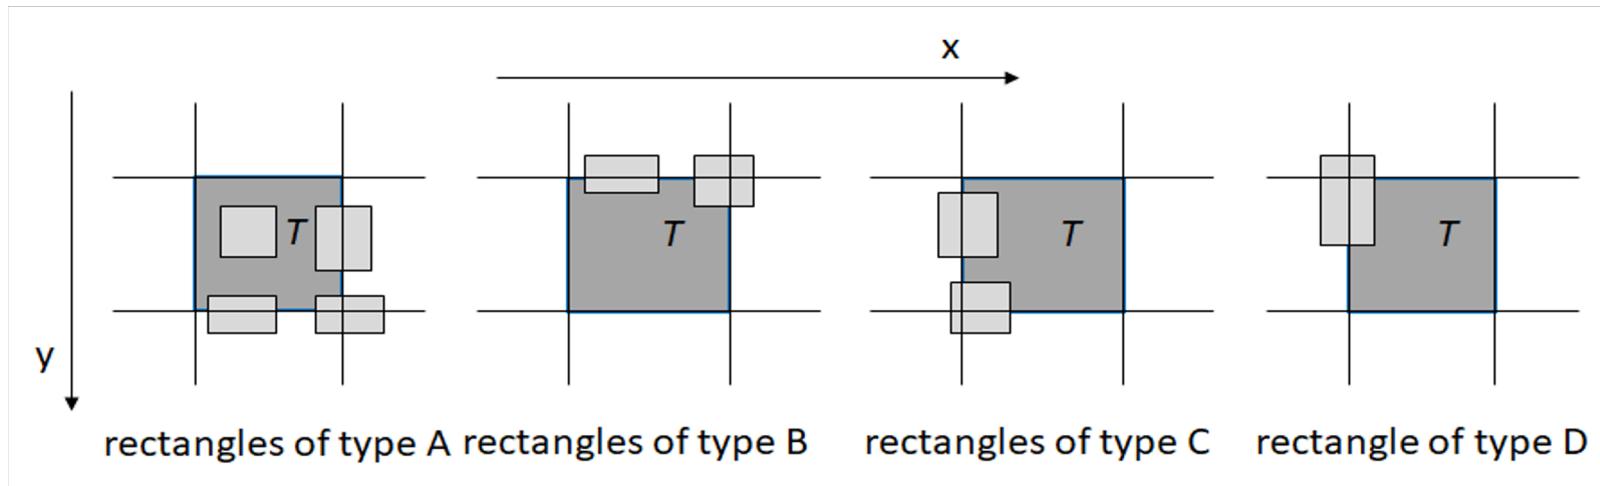
# Motivation (2/2)

- Core challenge
  - Duplicated results
- Previous techniques
  - Find all results then duplicate elimination
  - Hashing, sorting
  - Reference point (Dittrich & Seeger, ICDE 2000)
- Our technique
  - Duplicate avoidance
  - Divide each tile into four classes that cannot generate duplicates

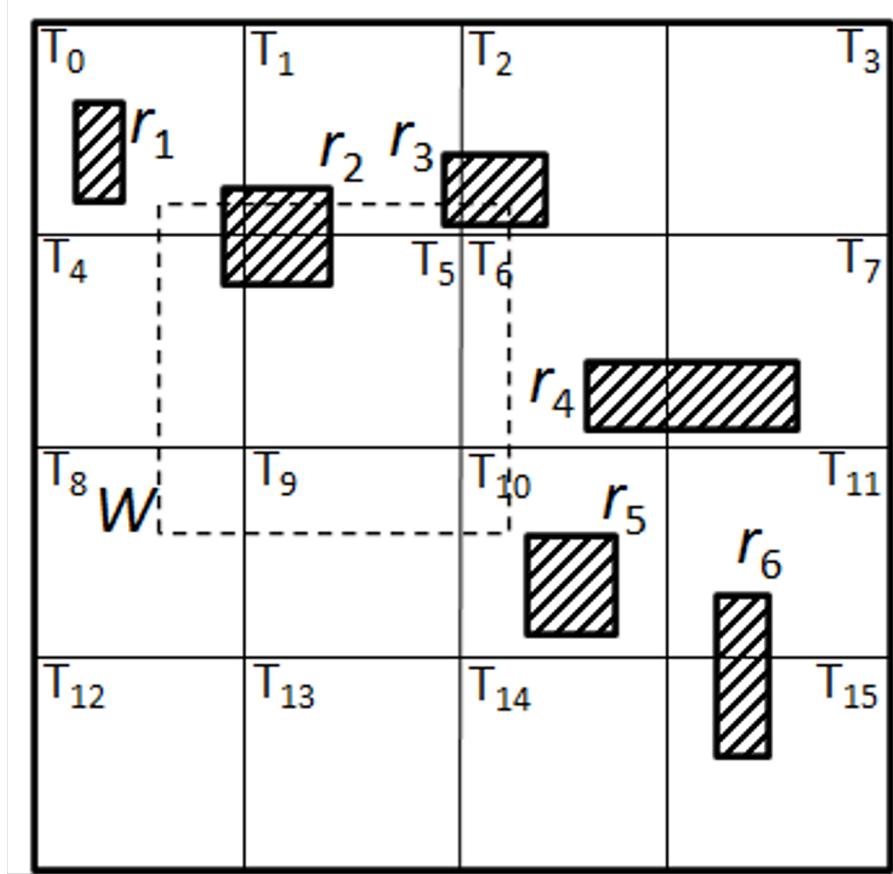


# Two-layer Partitioning

- First layer
  - Standard Space-Oriented Partitioning
  - Divide the space into disjoint spatial partitions, called tiles
- Second layer
  - Additional in our approach
  - Divide each tile into four classes A, B, C, and D



# Example



Tile	primary partitioning (standard)	secondary partitioning (our approach)
$T_0$	{ $r_1, r_2$ }	$A = \{r_1, r_2\}$
$T_1$	{ $r_2, r_3$ }	$A = \{r_3\}, C = \{r_2\}$
$T_2$	{ $r_3$ }	$C = \{r_3\}$
$T_4$	{ $r_2$ }	$B = \{r_2\}$
$T_5$	{ $r_2$ }	$D = \{r_2\}$
$T_6$	{ $r_4$ }	$A = \{r_4\}$
$T_7$	{ $r_4$ }	$C = \{r_4\}$
$T_{10}$	{ $r_5$ }	$A = \{r_5\}$
$T_{11}$	{ $r_6$ }	$A = \{r_6\}$
$T_{15}$	{ $r_6$ }	$B = \{r_6\}$

# Query processing

- Identify the tiles relevant to the query
- For each relevant tile
  - Select relevant classes
  - For each relevant class
    - Identify intersecting rectangles

# Query processing

- Identify the tiles relevant to the query
- For each relevant tile
  - Select relevant classes
  - For each relevant class
    - Identify intersecting rectangles

# Select Relevant Classes

		dimension x			
		T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>	T <sub>4</sub>
		A,B,C,D	A,B,C,D	A,B,C,D	A,B,C,D
		T <sub>5</sub>	T <sub>6</sub>	T <sub>7</sub>	T <sub>8</sub>
		A,B,C,D	A,B,C,D	A,B,C,D	A,B,C,D
		T <sub>9</sub>	T <sub>10</sub>	T <sub>11</sub>	T <sub>12</sub>
		A,B,C,D	A,B,C,D	A,B,C,D	A,B,C,D

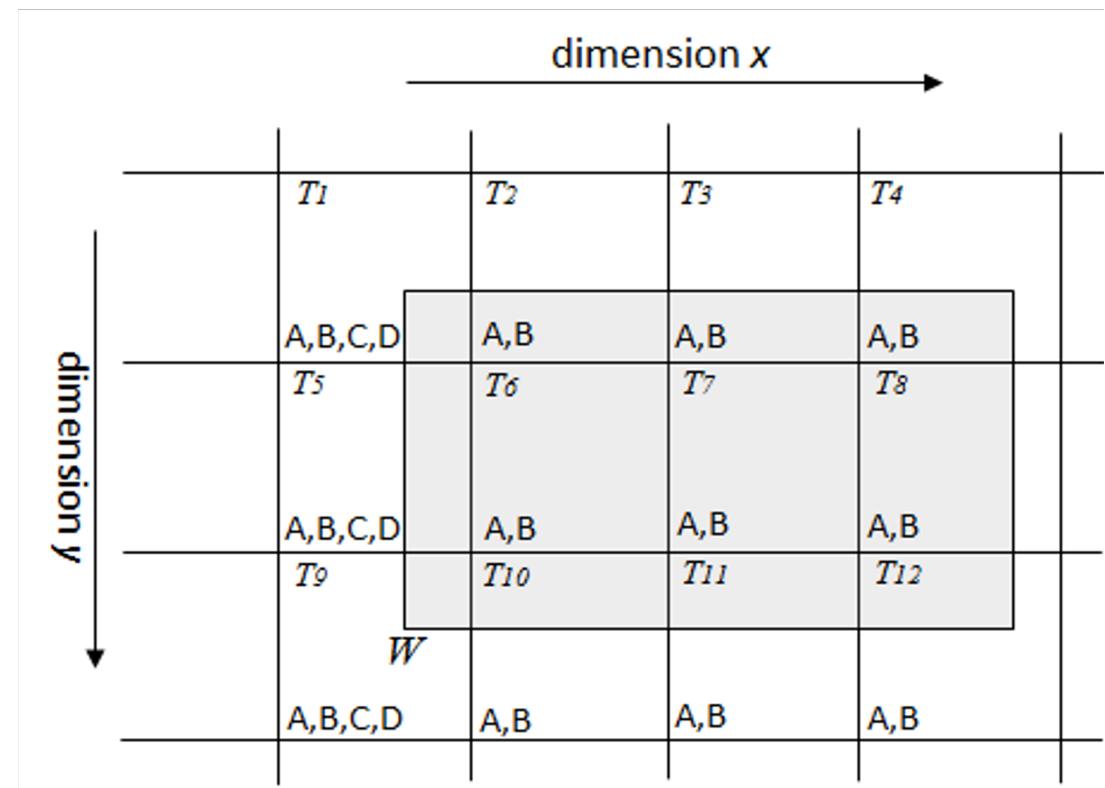
dimension y ↓

W

# Select Relevant Classes

- **Observation 1**

- Query W **intersects** tile T
- W **starts before** T in dimension x
- **Disregard** classes C and D
- Examples: T2, T3, T4, T6, T7, T8, T10, T11, T12



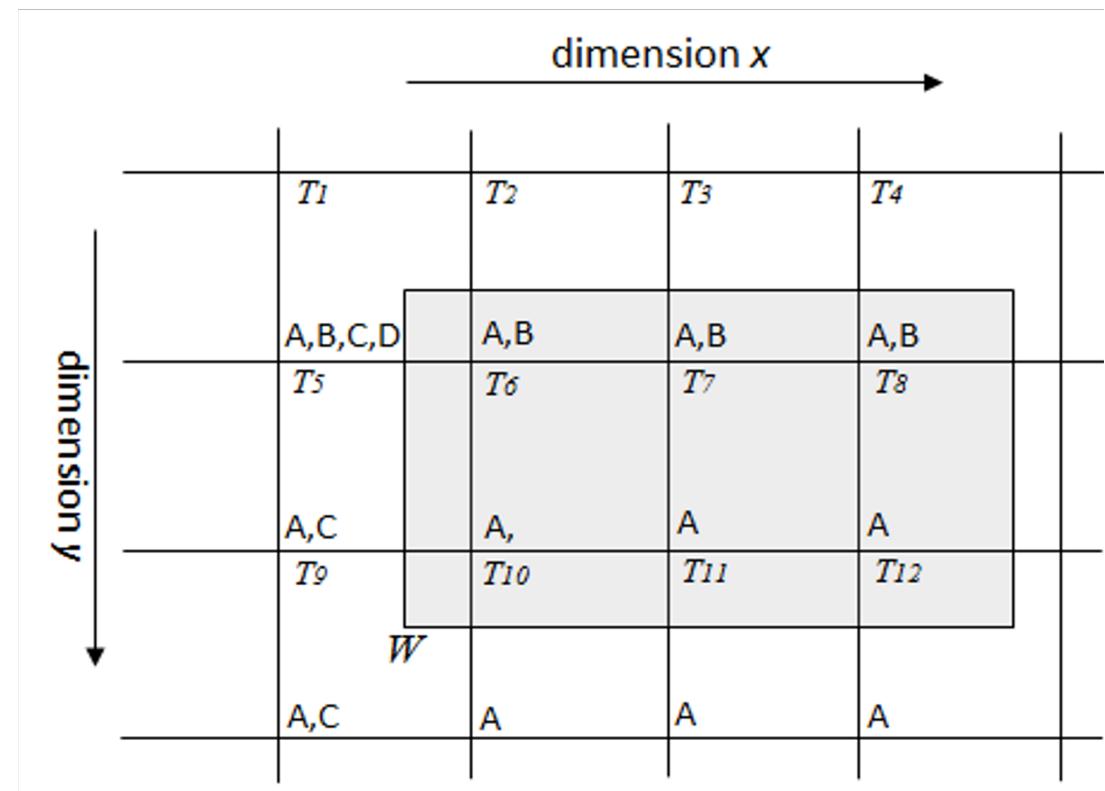
# Select Relevant Classes

- **Observation 1**

- Query W **intersects** tile T
- W **starts before** T in dimension x
- **Disregard** classes C and D
- Examples: T2, T3, T4, T6, T7, T8, T10, T11, T12

- **Observation 2**

- Query W **intersects** tile T
- W **starts before** T in dimension y
- **Disregard** classes B and D
- Examples: T5, T6, T7, T8, T9, T10, T11, T12



# Identify Intersecting Rectangles

- **Intersection test**
  - Normally, four comparisons
  - Minimize the number of comparisons if W is bigger than tile size
- **Observation 1**
  - Tile covered by window in a dimension
  - No intersection test in this dimension
- **Observation 2**
  - Query W ends in a Tile T
  - W starts before T in dimension d
  - One comparison : rectangle.dl  $\leq$  W.du
- **Observation 3**
  - Query W starts in a tile T
  - W ends after T in dimension d
  - One comparison: rectangle.du  $\geq$  W.dl

		dimension x						
		T1	T2	T3	T4			
dimension y	A,B,C,D	$r.x_u \geq W.x_l$ $r.y_u \geq W.y_l$	$r.y_u \geq W.y_l$	$r.y_u \geq W.y_l$	$r.y_u \geq W.y_l$ $r.x_l \leq W.x_u$			
	A,C	$r.x_u \geq W.x_l$	A	T6	A	T7	A	T8
dimension y	A,C	$r.x_u \geq W.x_l$ $r.y_l \leq W.y_u$	A	T10	A	T11	A	T12
	W	$r.y_l \leq W.y_u$	$r.y_l \leq W.y_u$	$r.y_l \leq W.y_u$	$r.y_l \leq W.y_u$	$r.x_l \leq W.x_u$	$r.y_l \leq W.y_u$	

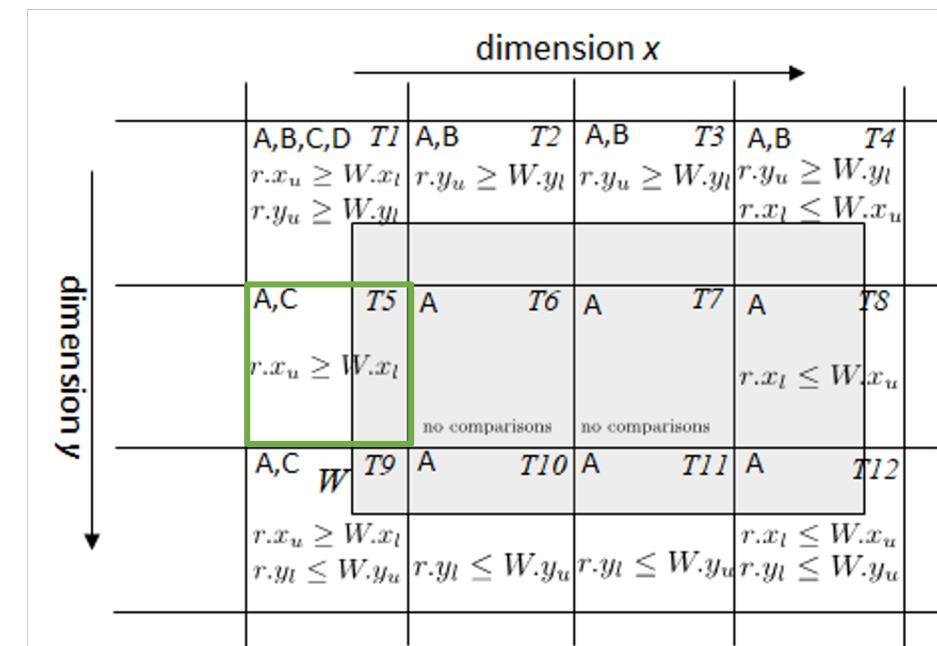
# Identify Intersecting Rectangles

- **Intersection test**
  - Normally, four comparisons
  - Minimize the number of comparisons if W is bigger than tile size
- **Observation 1**
  - Tile covered by window in a dimension
  - No intersection test in this dimension
- **Observation 2**
  - Query W ends in a Tile T
  - W starts before T in dimension d
  - One comparison : rectangle.dl  $\leq$  W.du
- **Observation 3**
  - Query W starts in a tile T
  - W ends after T in dimension d
  - One comparison: rectangle.du  $\geq$  W.dl

		dimension x					
		T1	T2	T3	T4		
dimension y	A,B,C,D	$r.x_u \geq W.x_l$	$r.y_u \geq W.y_l$	$r.y_u \geq W.y_l$	$r.y_u \geq W.y_l$		
		$r.y_u \geq W.y_l$				$r.x_l \leq W.x_u$	
		T5	A	T6	A	T7	A
		$r.x_u \geq W.x_l$				$r.x_l \leq W.x_u$	
		A,C	$r.y_u \geq W.y_l$			no comparisons	no comparisons
		$r.x_u \geq W.x_l$					
		A,C	T9	A	T10	A	T11
		$r.x_u \geq W.x_l$	$r.y_l \leq W.y_u$	$r.y_l \leq W.y_u$	$r.y_l \leq W.y_u$	$r.x_l \leq W.x_u$	$r.y_l \leq W.y_u$

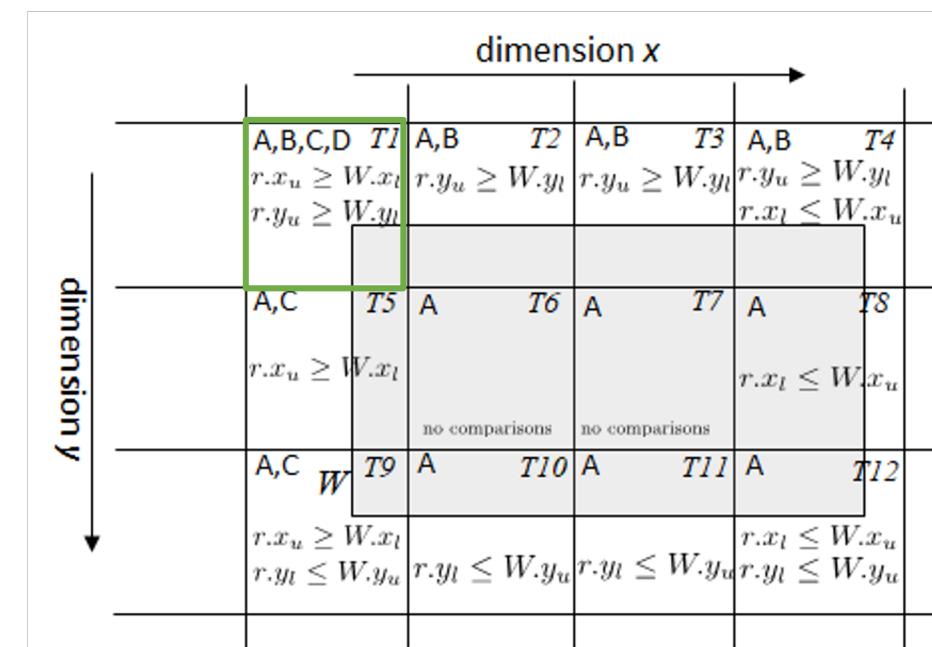
# Identify Intersecting Rectangles

- **Intersection test**
  - Normally, four comparisons
  - Minimize the number of comparisons if W is bigger than tile size
- **Observation 1**
  - Tile covered by window in a dimension
  - No intersection test in this dimension
- **Observation 2**
  - Query W ends in a Tile T
  - W starts before T in dimension d
  - One comparison : rectangle.dl  $\leq$  W.du
- **Observation 3**
  - Query W starts in a tile T
  - W ends after T in dimension d
  - One comparison: rectangle.du  $\geq$  W.dl



# Identify Intersecting Rectangles

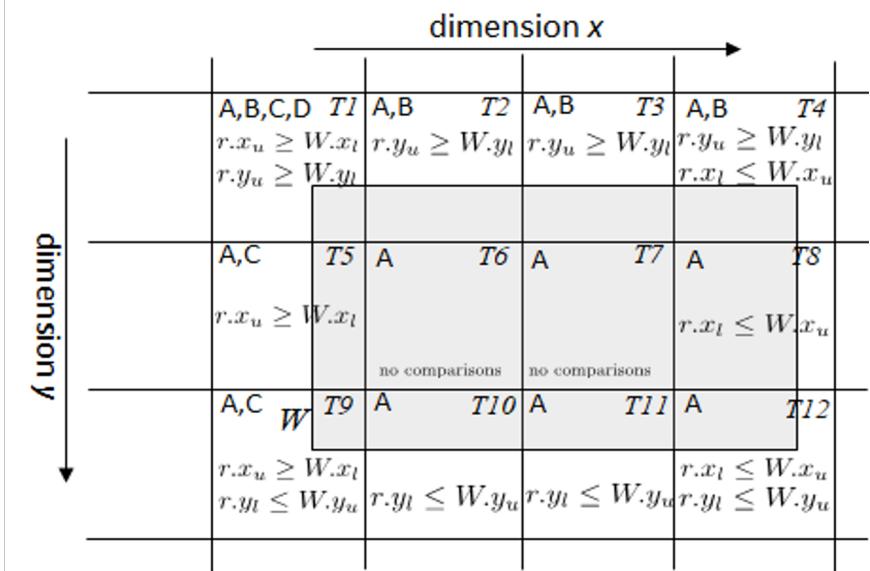
- **Intersection test**
  - Normally, four comparisons
  - Minimize the number of comparisons if W is bigger than tile size
- **Observation 1**
  - Tile covered by window in a dimension
  - No intersection test in this dimension
- **Observation 2**
  - Query W ends in a Tile T
  - W starts before T in dimension d
  - One comparison : rectangle.dl  $\leq$  W.du
- **Observation 3**
  - Query W starts in a tile T
  - W ends after T in dimension d
  - One comparison: rectangle.du  $\geq$  W.dl



# Storage Optimization

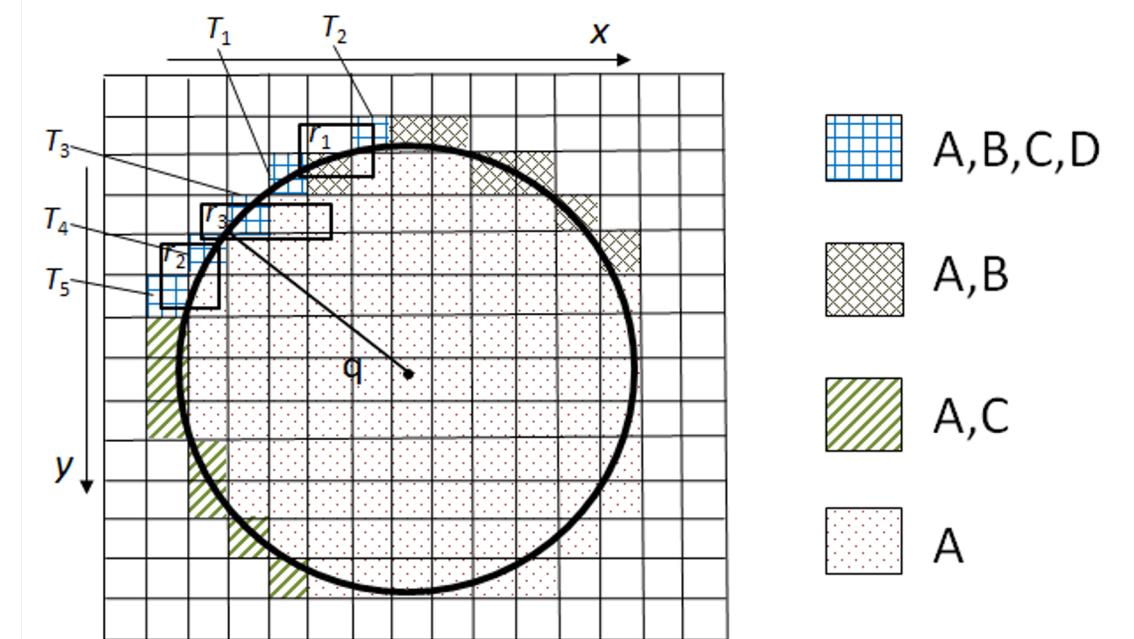
- Store the MBRs using the **Decomposition Storage Model**
  - Reduces the query cost
  - Improves data access locality
- Each table contains **(coordinate, id)** pairs
  - Sorted by coordinate
  - Used for queries where **one endpoint** is needed
- **Not necessary** to store all decompositions
- **Cons**
  - Requires **additional storage**
  - Expensive to **update**

partition	required tables
$T^A$	$L_{xl}^A, L_{xu}^A, L_{yl}^A, L_{yu}^A$
$T^B$	$L_{xl}^B, L_{xu}^B, L_{yu}^B$
$T^C$	$L_{xu}^C, L_{yl}^C, L_{yu}^C$
$T^D$	$L_{xu}^D, L_{yu}^D$



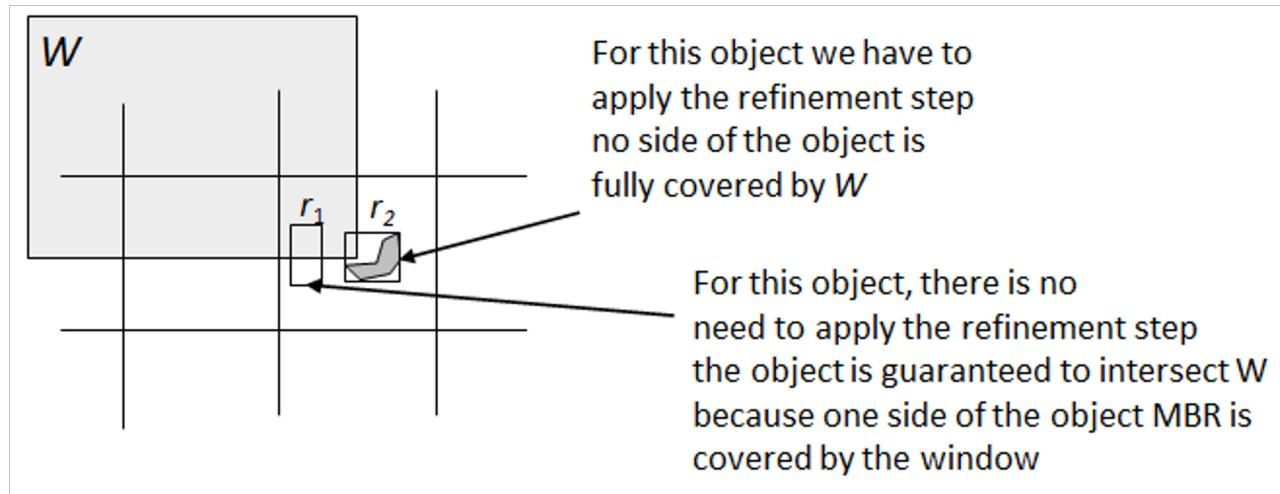
# Disk Query

- If previous tile in **x** dimension intersects with the query -> **A,C**
- If previous tile in **y** dimension intersects with the query -> **A,B**
- if tile is **full covered** by query -> **A**
- If previous tiles does **not intersect** with the query in both **dimensions** -> **A,B,C,D**



# Accelerate Refinement

- Secondary filtering before refinement
- Observation
  - At least one side of an MBR inside the query
  - Object does not need refinement step



# Batch Query Processing

- Single-threaded and multi-threaded implementations
- **Queries-based** approach
  - Process **each query** independently
  - **Assign queries** to available threads in round robin
  - Cache-agnostic
- **Tiles-based** approach
  - For each tile
    - Find intersecting queries
    - Combine relevant classes for all queries
  - Process **each tile** independently
  - **Assign tiles** to available threads in round robin
  - Cache-conscious

# Setup & Datasets

- Hardware:
  - Processor: dual Index(R) Xeon(R) CPU E5-2630 v4 clocked at 2.20Ghz with 384 GBs of RAM
  - Hyper-threading enabled for batch processing, up to 40 threads
- Implementation:
  - Programming Language: C++
  - Operation System: CentOs Linux 7.6.1810

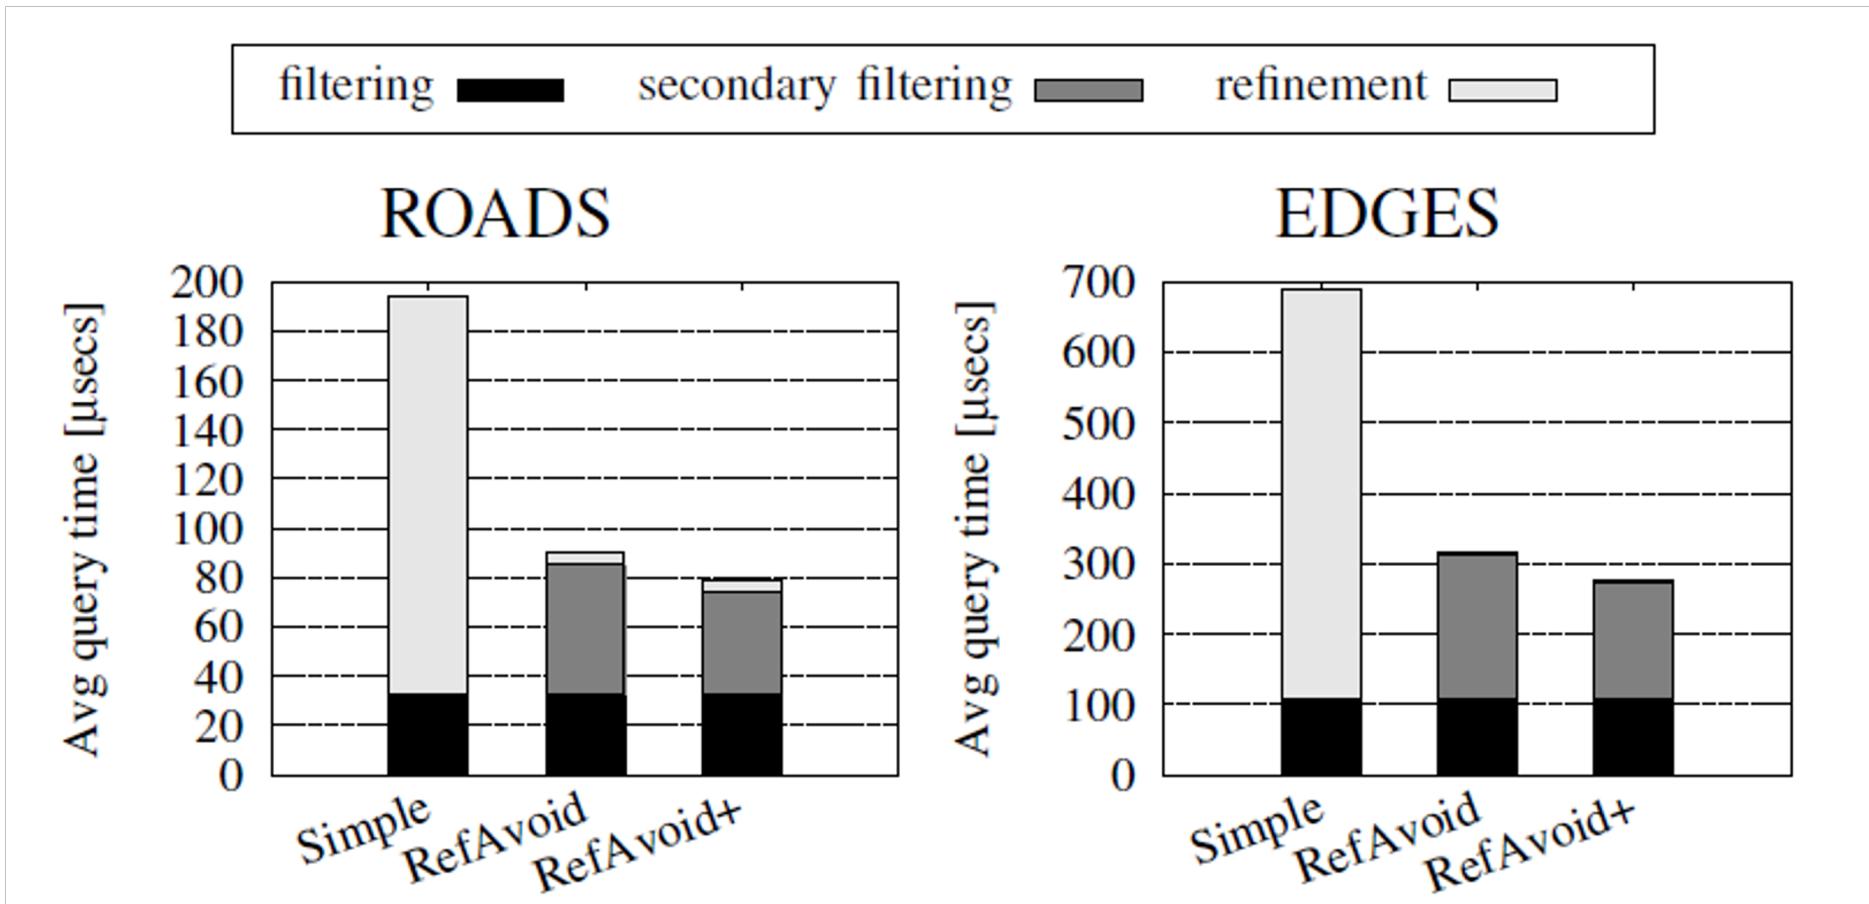
Synthetic dataset

parameter	values	Default
cardinality	1M, 5M, 10M, 50M, 100M	10M
area	$10^{-\infty}$ , $10^{-14}$ , $10^{-12}$ , $10^{-8}$ , $10^{-6}$	$10^{-10}$
distribution	Uniform or Zipfian ( $a = 1$ )	-

Real dataset

dataset	type	card.	avg. x-extend	avg. y-extend
ROADS	linestrings	20M	0.00001173	0.00000915
EDGES	polygons	70M	0.00000491	0.00000383
TIGER	mixed	98M	0.00000740	0.00000576

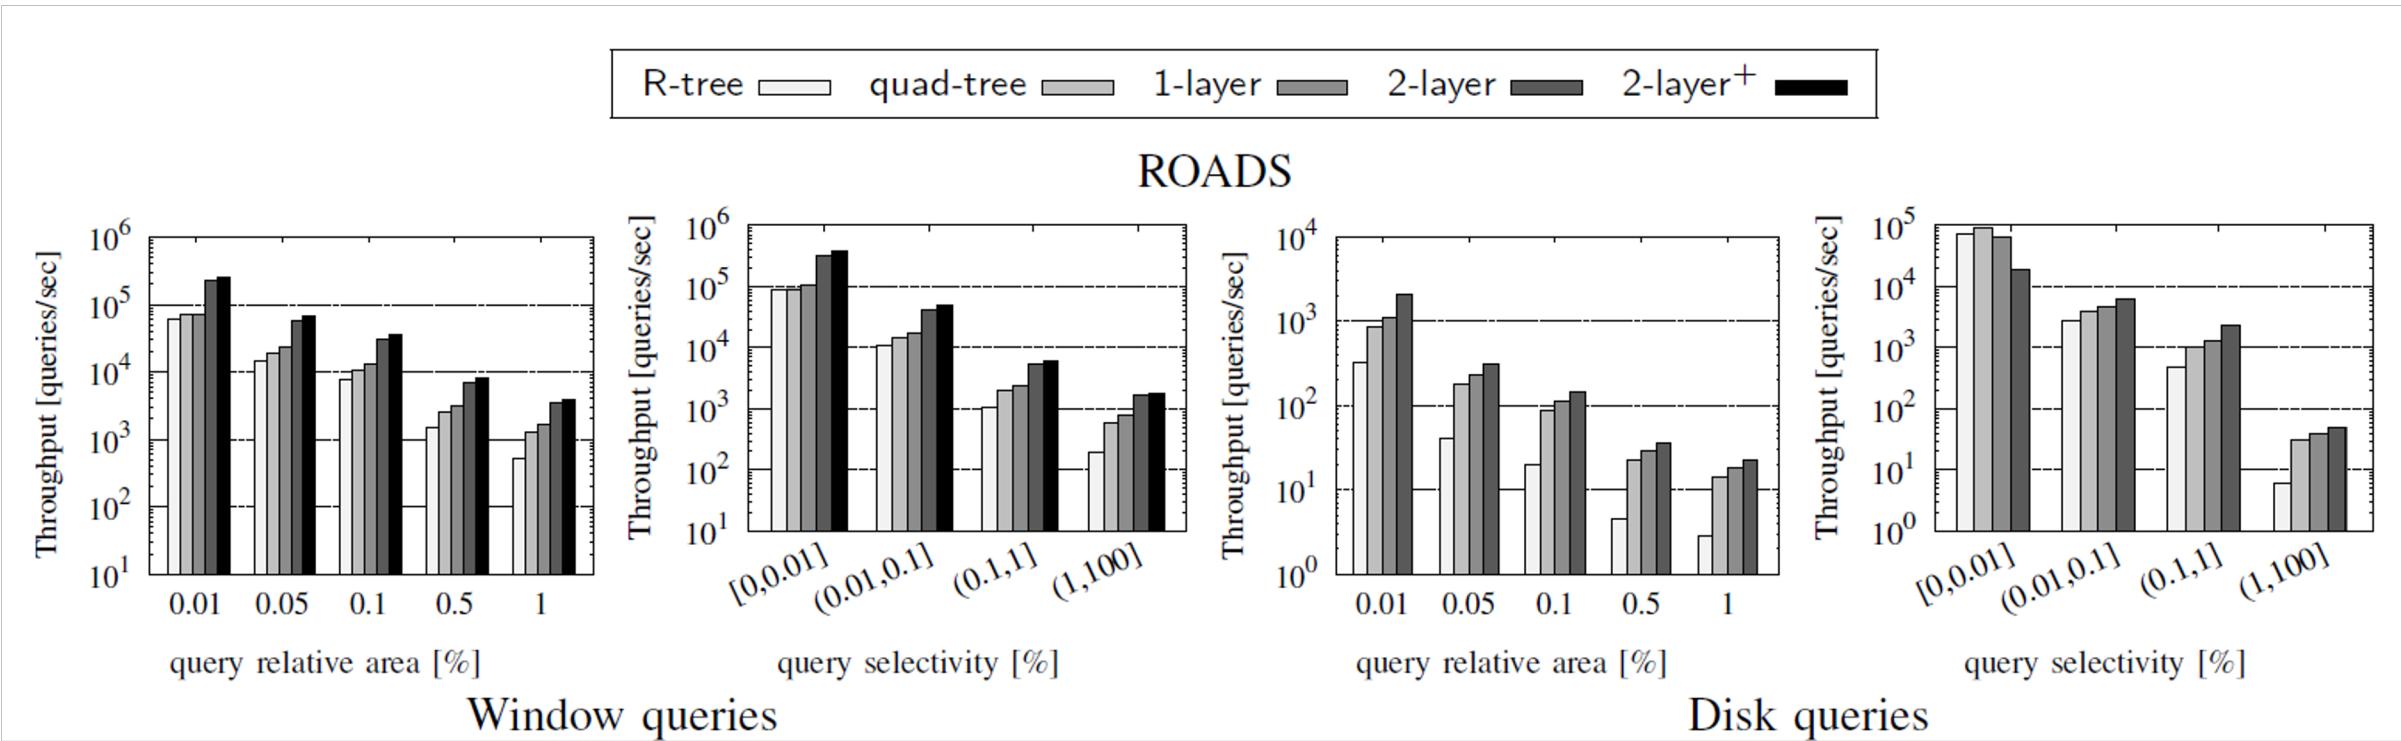
# Filtering vs Refinement



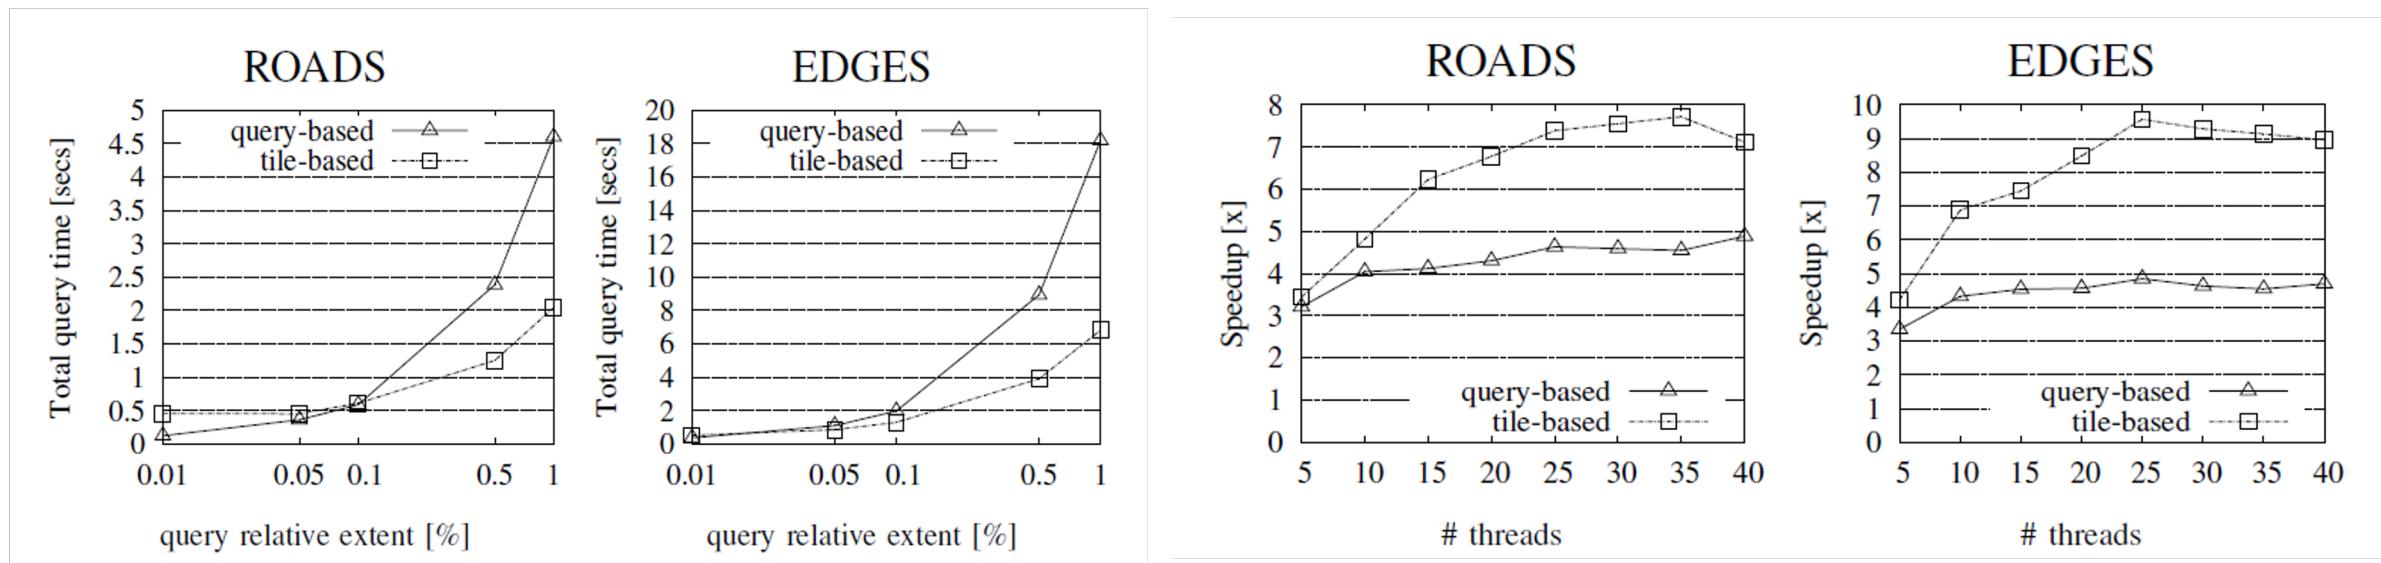
# Compared Methods and their Throughput

Type	index	throughput (queries/sec)	
		ROADS	EDGES
SOP	2-layer	30981	9406
	2-layer <sup>+</sup>	36444	10855
	1-layer	12597	4403
	quad-tree	10949	3640
	quad-tree, 2-layer	16883	5831
DOP	R-tree	7888	2011
	R*-tree	6415	1610
	BLOCK	< 1	< 1
	MXCIF quad-tree	8	2

# Query Processing: Real Data



# Batch query processing (window queries)



# Conclusion

- Contributions
  - Our **two-layer** partitioning:
    - Easy to implement
    - Can be applied to any SOP index
    - Reduces the number of comparisons
    - Duplicate avoidance
  - Our secondary filtering technique avoids refinement step for the majority of the query results
  - Efficient processing of multiple queries in batch and in parallel
- Future work
  - Implementation of two-layer partitioning for 3D data
  - Apply two-layer partitioning in distributed data management
  - Consider other popular queries types, such as KNN and spatial join

Thank you for your attention

Questions

?