

# A Forward Scan based Plane Sweep Algorithm for Parallel Interval Joins

Panagiotis Bouros<sup>1</sup> and Nikos Mamoulis<sup>2</sup>

<sup>1</sup> Aarhus University, Denmark

<sup>2</sup> University of Ioannina, Greece

# Interval Joins

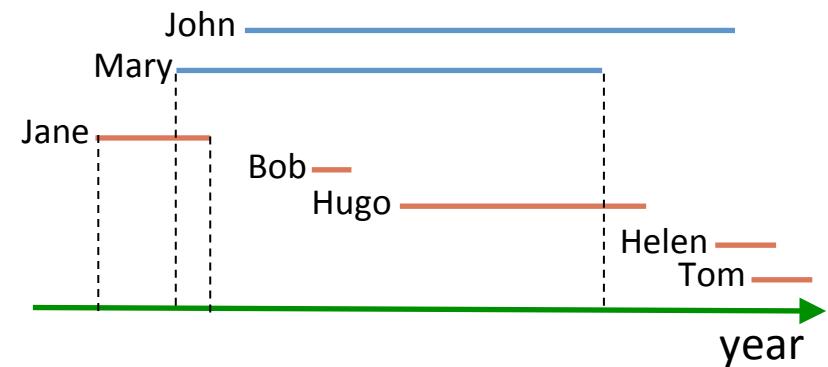
employee	start	end
John	1994	2006
Mary	1992	2002

employee	start	end
Jane	1990	1993
Bob	1995	1996
Hugo	1997	2003
Helen	2005	2007
Tom	2006	2008

# Interval Joins

employee	start	end
John	1994	2006
Mary	1992	2002

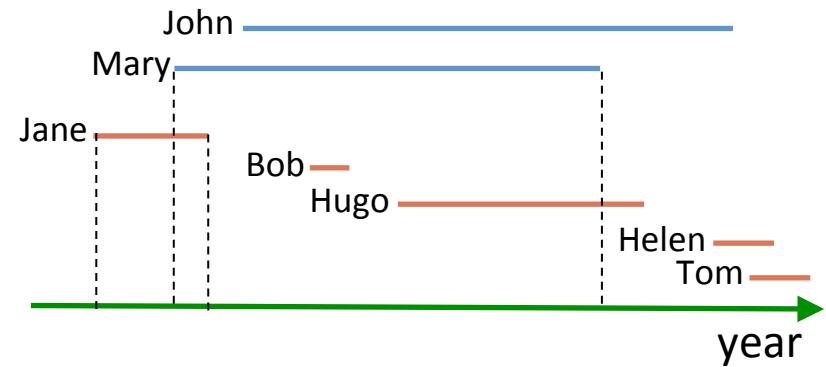
employee	start	end
Jane	1990	1993
Bob	1995	1996
Hugo	1997	2003
Helen	2005	2007
Tom	2006	2008



# Interval Joins

employee	start	end
John	1994	2006
Mary	1992	2002

employee	start	end
Jane	1990	1993
Bob	1995	1996
Hugo	1997	2003
Helen	2005	2007
Tom	2006	2008

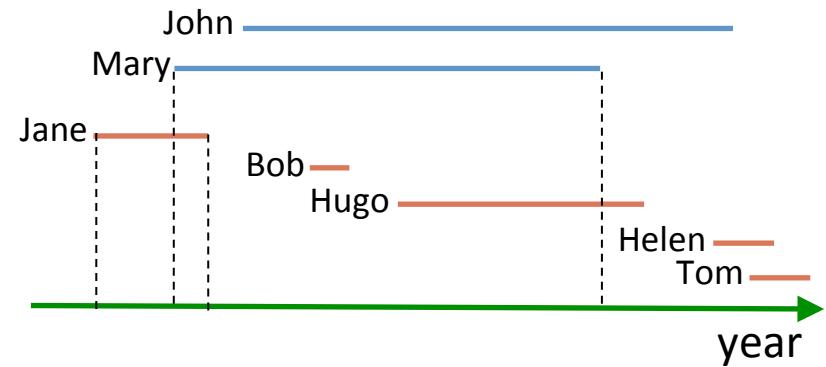


Find all pair of employees whose period of work on D1 and D2 intersect

# Interval Joins

employee	start	end
John	1994	2006
Mary	1992	2002

employee	start	end
Jane	1990	1993
Bob	1995	1996
Hugo	1997	2003
Helen	2005	2007
Tom	2006	2008



Find all pair of employees whose **period of work** on **D1** and **D2** **intersect**

- **Applications**
  - Temporal databases
  - Multidimensional data management
  - Uncertain data management

# Our focus

- Efficient evaluation of interval joins
  - Single-threaded processing
    - Simple plane sweep based method
    - Competitive to state-of-the-art
  - Parallel processing
    - Partitioning-based join
    - Share nothing

# **SINGLE-THREADED PROCESSING**

# Related work

## Nested loops & Sort-merge join

[Segev and Gunadhi, VLDB'89]  
[Gunadhi and Segev, ICDE'91]

## Index-based

[Zhang et al., ICDE'02]  
[Enderle et al., SIGMOD'04]

## Partitioning-based

[Soo et al., ICDE'94]  
[Dignös et al., SIGMOD'14]  
[Cafagna and Böhlen, VLDBJ'17]

## Plane-sweep based

[Brinkhoff et al., SIGMOD'93]  
[Arge et al., VLDB'98]  
[Piatov et al., ICDE'16]

# Related work

## Nested loops & Sort-merge join

[Segev and Gunadhi, VLDB'89]  
[Gunadhi and Segev, ICDE'91]

## Index-based

[Zhang et al., ICDE'02]  
[Enderle et al., SIGMOD'04]

## Partitioning-based

[Soo et al., ICDE'94]  
[Dignös et al., SIGMOD'14]  
[Cafagna and Böhlen, VLDBJ'17]

## Plane-sweep based

[Brinkhoff et al., SIGMOD'93]  
[Arge et al., VLDB'98]  
[Piatov et al., ICDE'16]

# Plane sweep methods

- Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]

- Sweep line stops both on *start* and *end*
- Backwards scan, Gapless hash map buffers open intervals

## Pros

- ✓ No domain-point comparisons
- ✓ Tailored to modern hardware
- ✓ Main memory cache-aware
- ✓ Fast

## Cons

- ✗ Special structure needed

# Plane sweep methods

- Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]

- Sweep line stops both on *start* and *end*
- Backwards scan, Gapless hash map buffers open intervals

**Pros**

- ✓ No domain-point comparisons
- ✓ Tailored to modern hardware
- ✓ Main memory cache-aware
- ✓ Fast

**Cons**

- ✗ Special structure needed

- Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]

- Sweep lines stops only on *start*

**Pros**

- ✓ Simple
- ✓ No special structure needed

**Cons**

- ✗  $|R| + |S| + |R \bowtie S|$  comparisons in total

# Plane sweep methods

- Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]

- Sweep line stops both on *start* and *end*
- Backwards scan, Gapless hash map buffers open intervals

**Pros**

- ✓ No domain-point comparisons
- ✓ Tailored to modern hardware
- ✓ Main memory cache-aware
- ✓ Fast

**Cons**

- ✗ Special structure needed

- Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]

- Sweep lines stops only on *start*

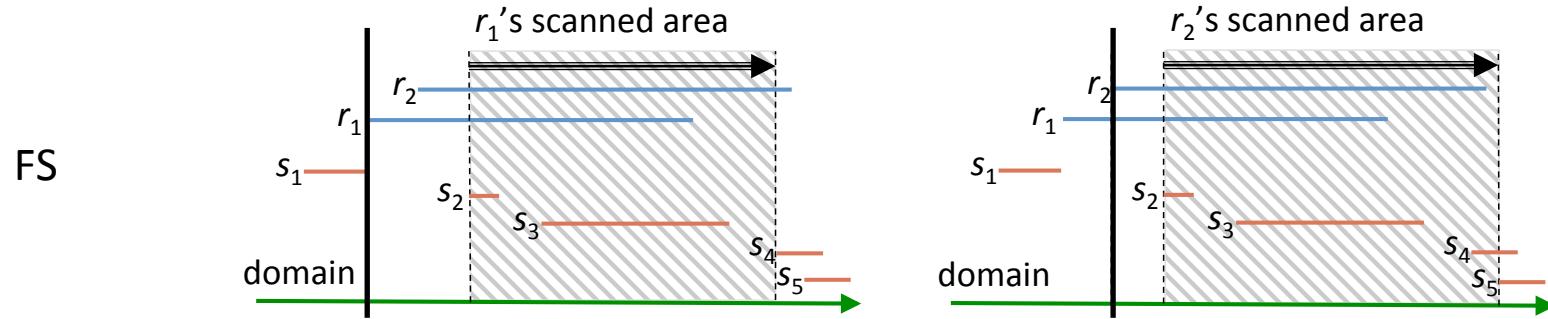
**Pros**

- ✓ Simple
- ✓ No special structure needed

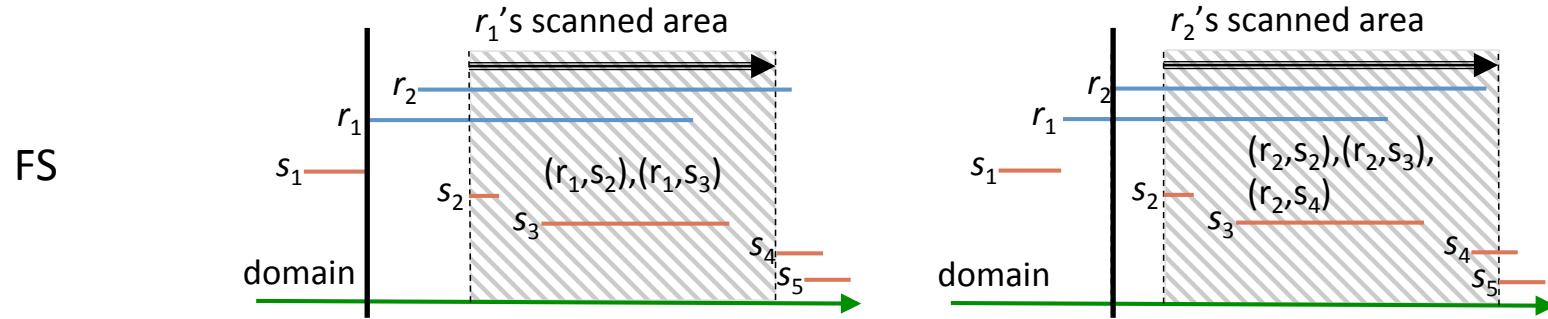
**Cons**

- ✗  $|R| + |S| + |R \times S|$  comparisons in total

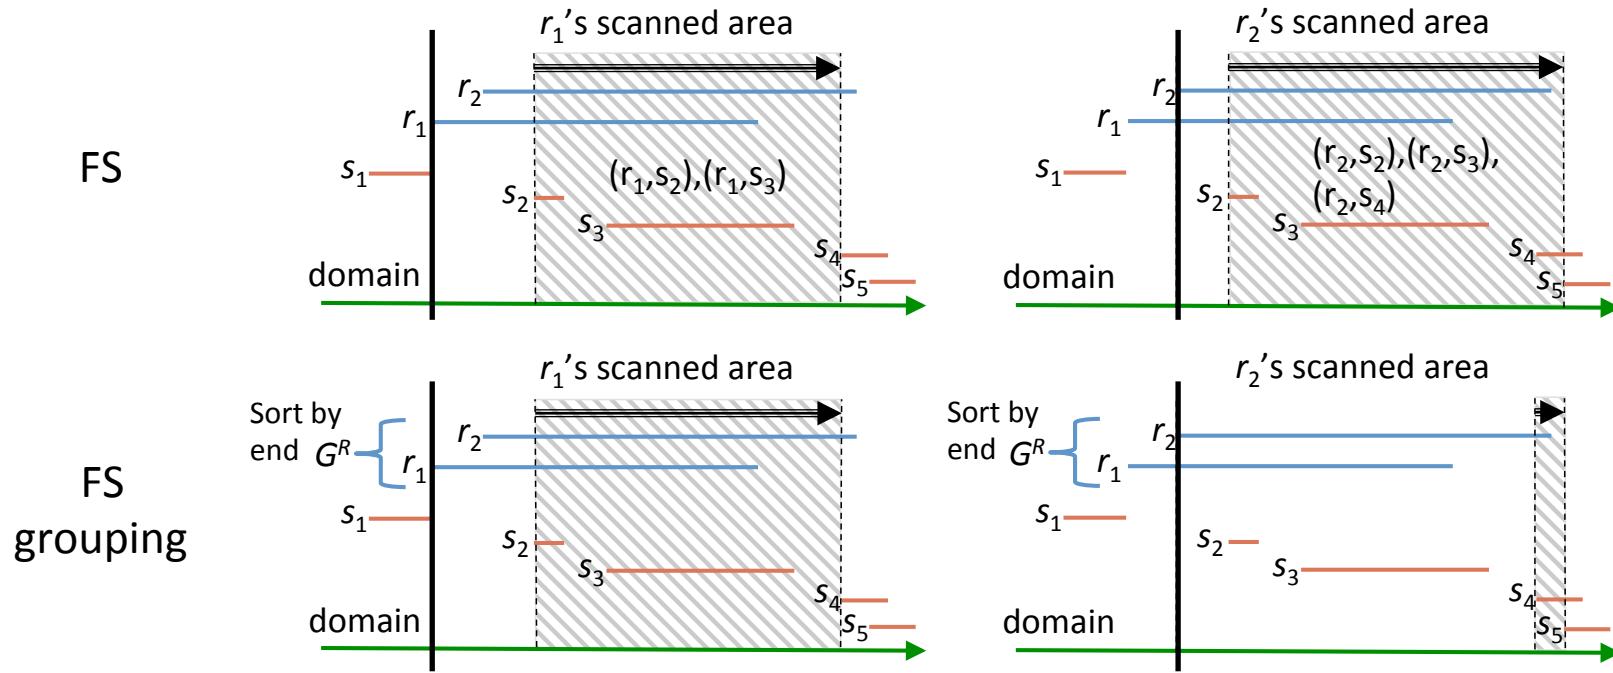
# Optimizing FS



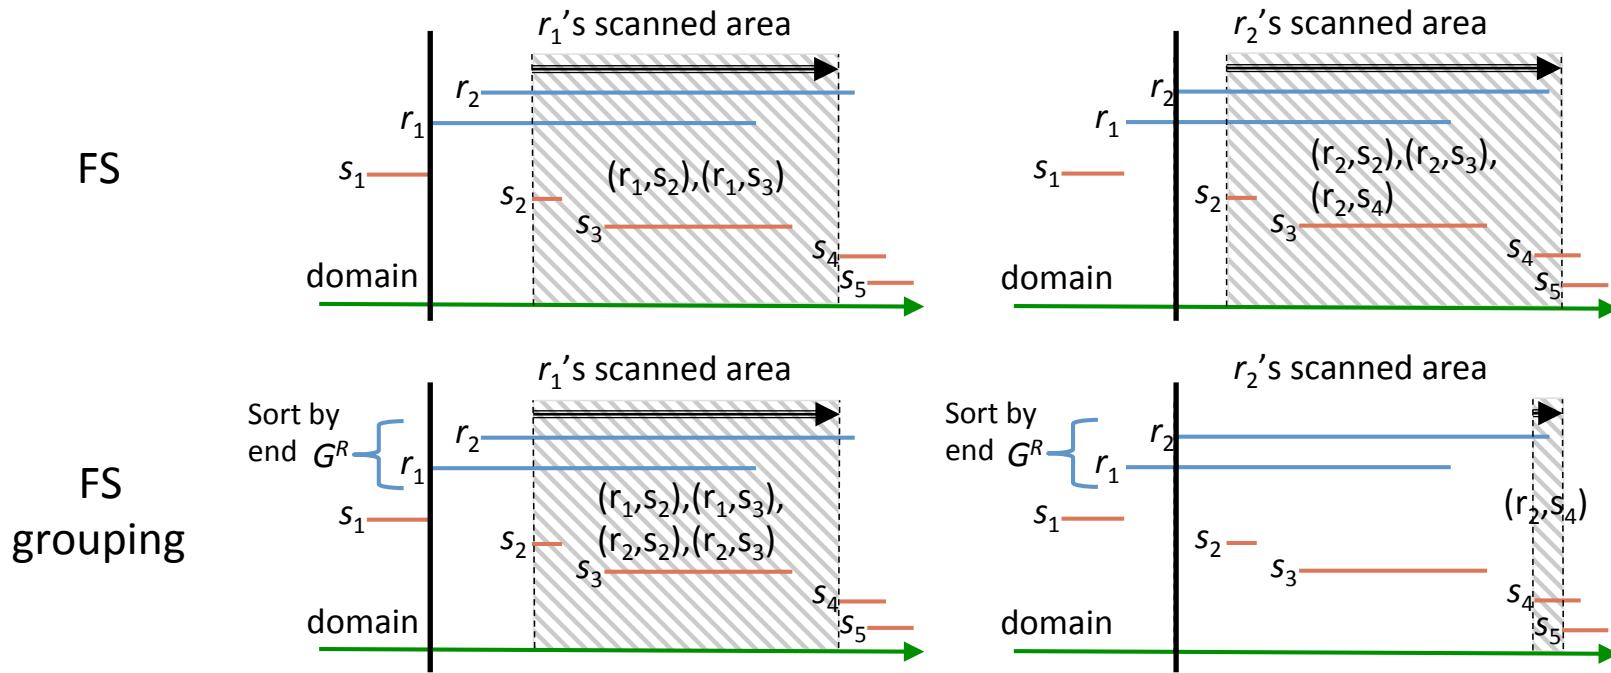
# Optimizing FS



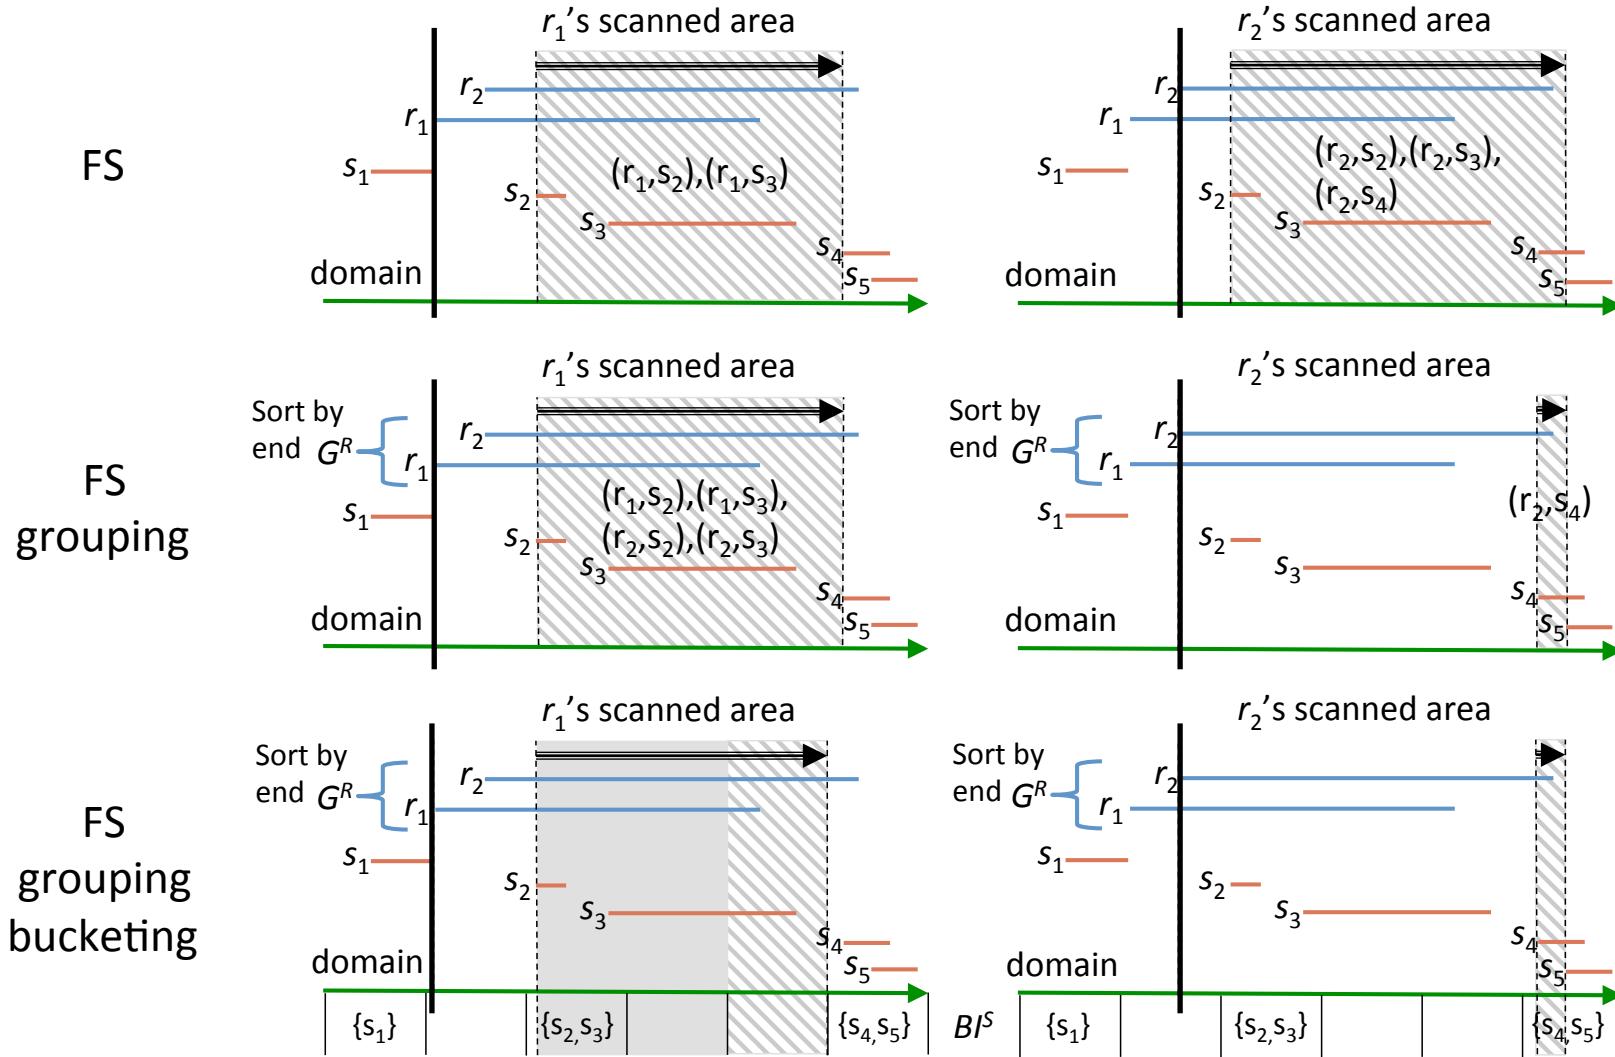
# Optimizing FS



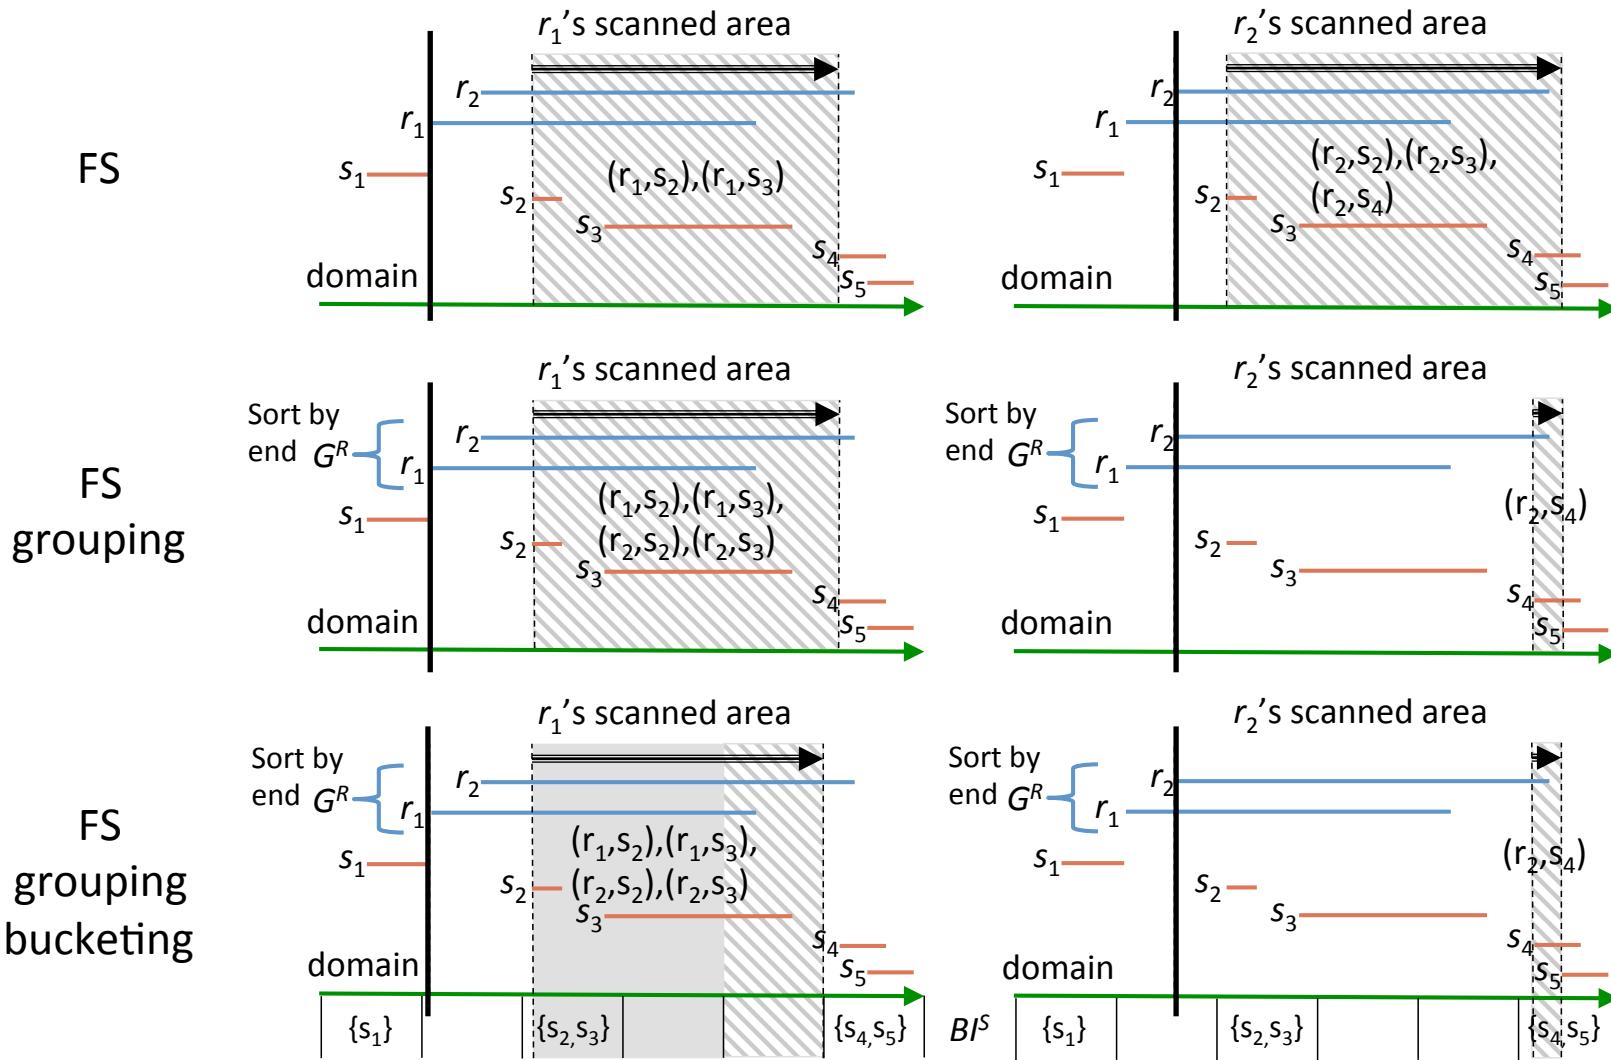
# Optimizing FS



# Optimizing FS



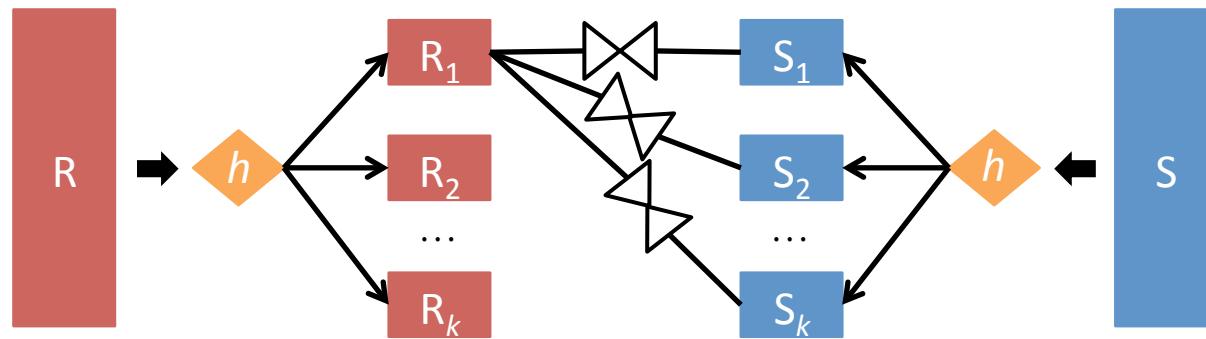
# Optimizing FS



# **PARALLEL PROCESSING**

# Hash-based partitioning

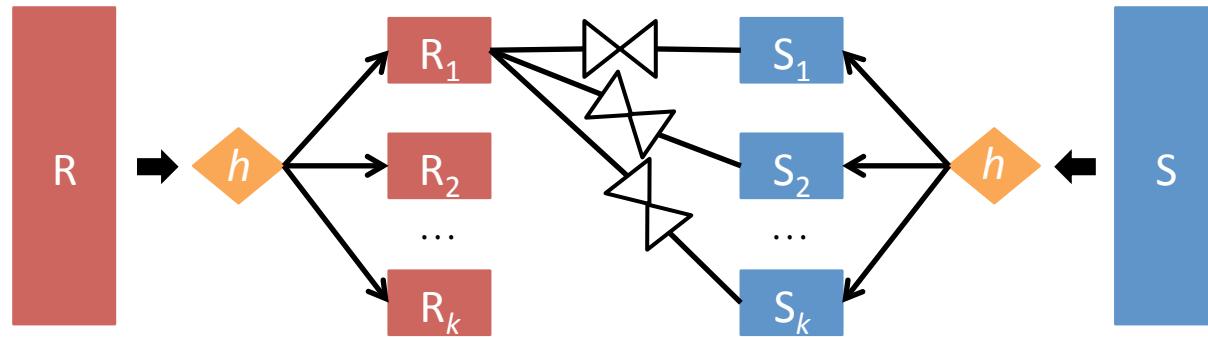
[Piatov et al., ICDE'16]



- Idea
  - Randomly split each input into  $k$  partitions using hash function  $h$
  - Evaluate  $k^2$  independent partition joins

# Hash-based partitioning

[Piatov et al., ICDE'16]



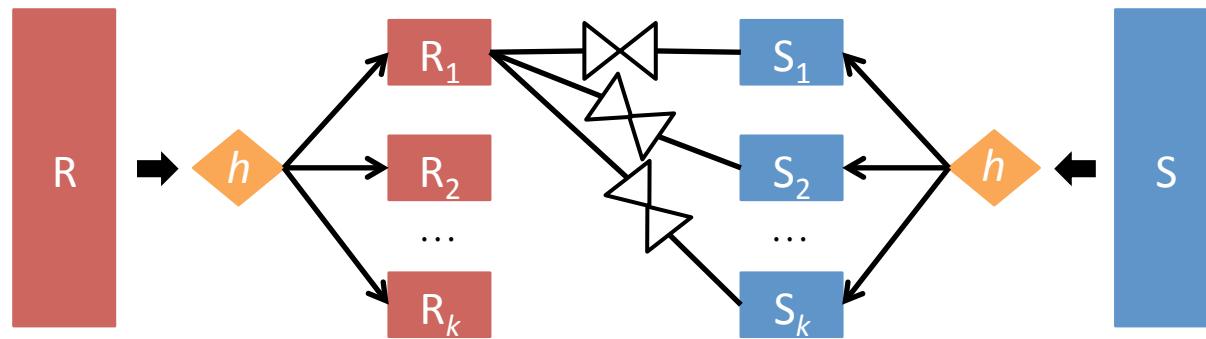
- Idea
  - Randomly split each input into  $k$  partitions using hash function  $h$
  - Evaluate  $k^2$  independent partition joins

## Pros

- ✓ Simple
- ✓ Load balancing

# Hash-based partitioning

[Piatov et al., ICDE'16]



- **Idea**
  - Randomly split each input into  $k$  partitions using hash function  $h$
  - Evaluate  $k^2$  independent partition joins

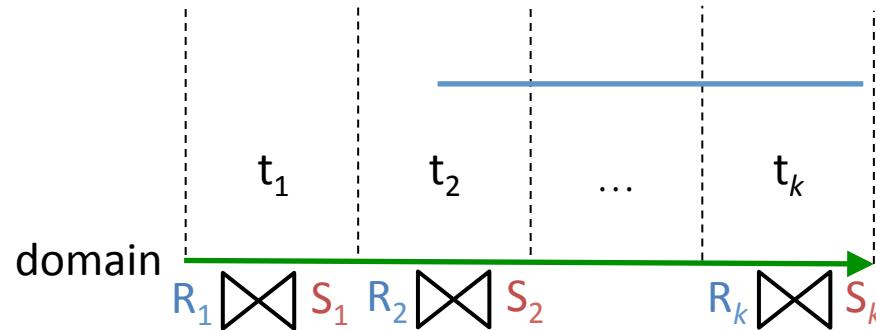
## Pros

- ✓ Simple
- ✓ Load balancing

## Cons

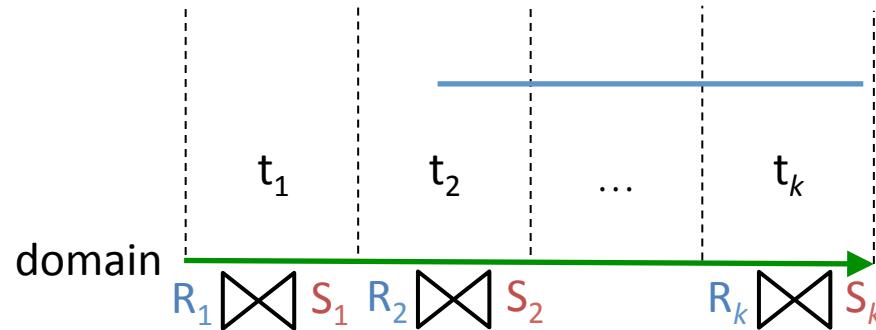
- ✗ Domain-point comparisons rise
  - $2 \cdot k \cdot (|R| + |S|)$  for EBI/LEBI,  $k \cdot (|R| + |S|)$  for FS
- ✗ Degree of parallelism
  - $n$  CPU cores  $\rightarrow k = \sqrt{n}$  partitions

# Domain-based partitioning



- **Idea**
  - Split domain into  $k$  tiles
  - Replicate intervals
  - Evaluate  $k$  independent partition joins

# Domain-based partitioning

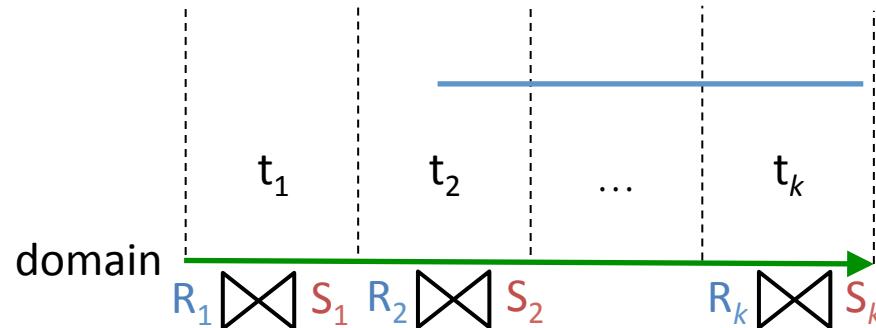


- **Idea**
  - Split domain into  $k$  tiles
  - Replicate intervals
  - Evaluate  $k$  independent partition joins

## Pros

- ✓ Degree of parallelism
  - $n$  CPU cores  $\rightarrow k = n$  partitions
- ✓ Automatic duplicate elimination

# Domain-based partitioning



- **Idea**
  - Split domain into  $k$  tiles
  - Replicate intervals
  - Evaluate  $k$  independent partition joins

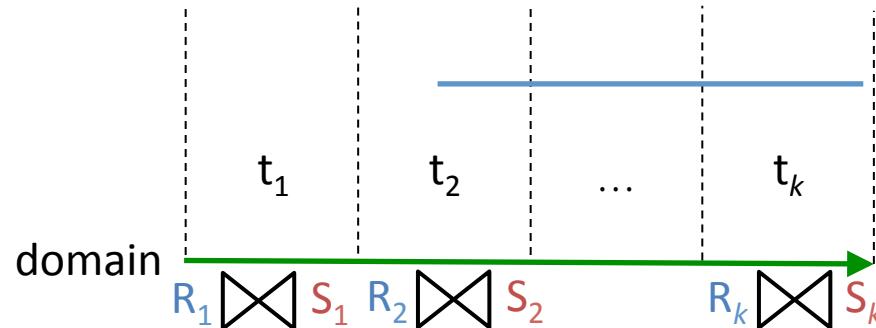
## Pros

- ✓ Degree of parallelism
  - $n$  CPU cores  $\rightarrow k = n$  partitions
- ✓ Automatic duplicate elimination

## Cons

- ✗ Replication
- ✗ Load balancing

# Domain-based partitioning



- **Idea**
  - Split domain into  $k$  tiles
  - Replicate intervals
  - Evaluate  $k$  independent partition joins

## Pros

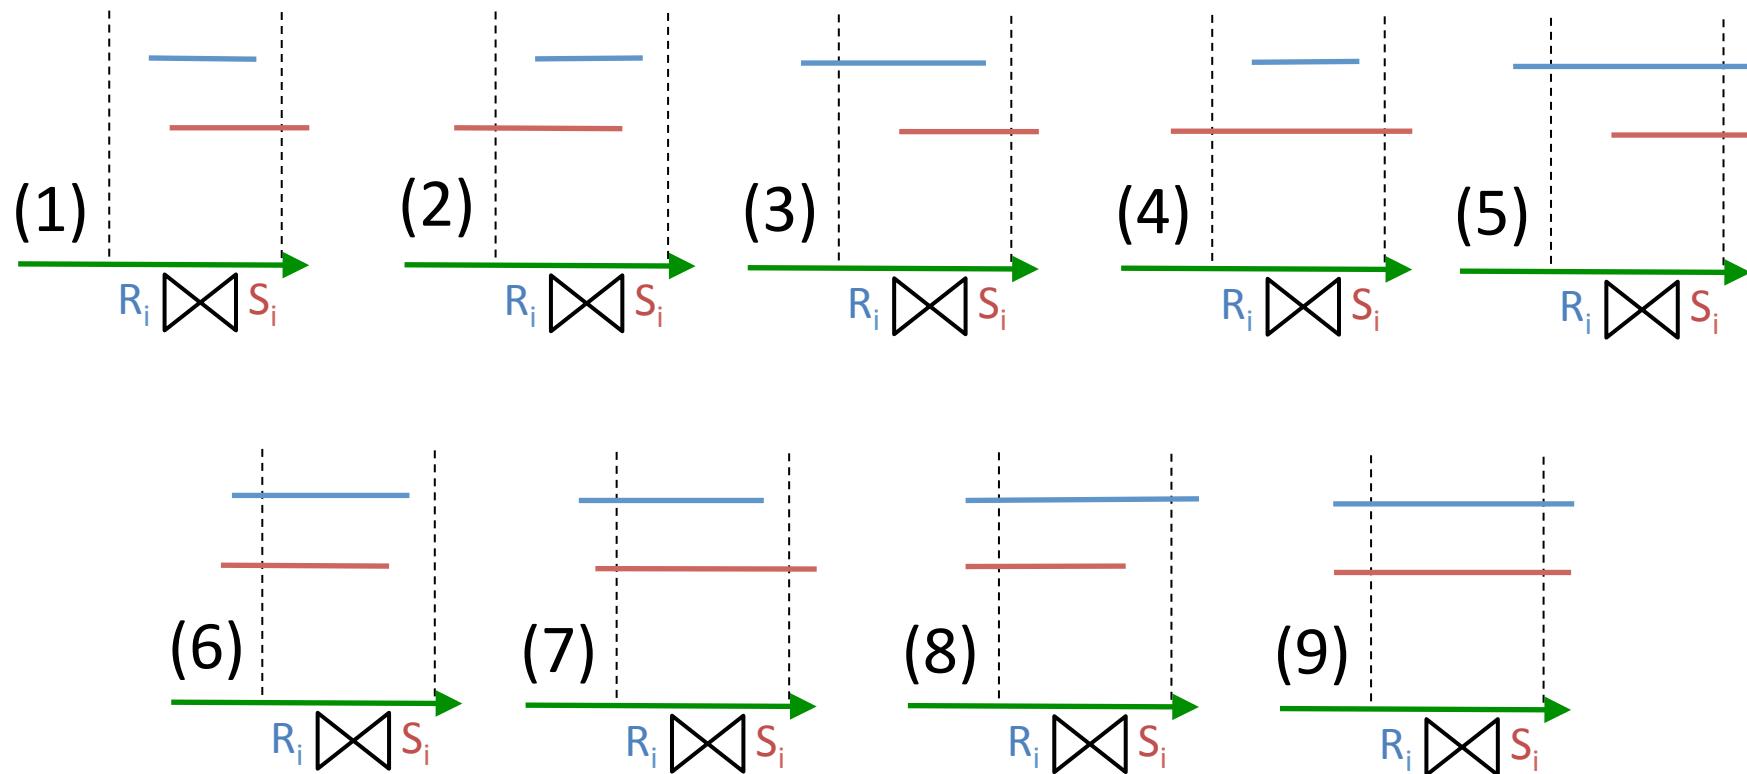
- ✓ Degree of parallelism
  - $n$  CPU cores  $\rightarrow k = n$  partitions
- ✓ Automatic duplicate elimination

## Cons

- ✗ Replication
- ✗ Load balancing

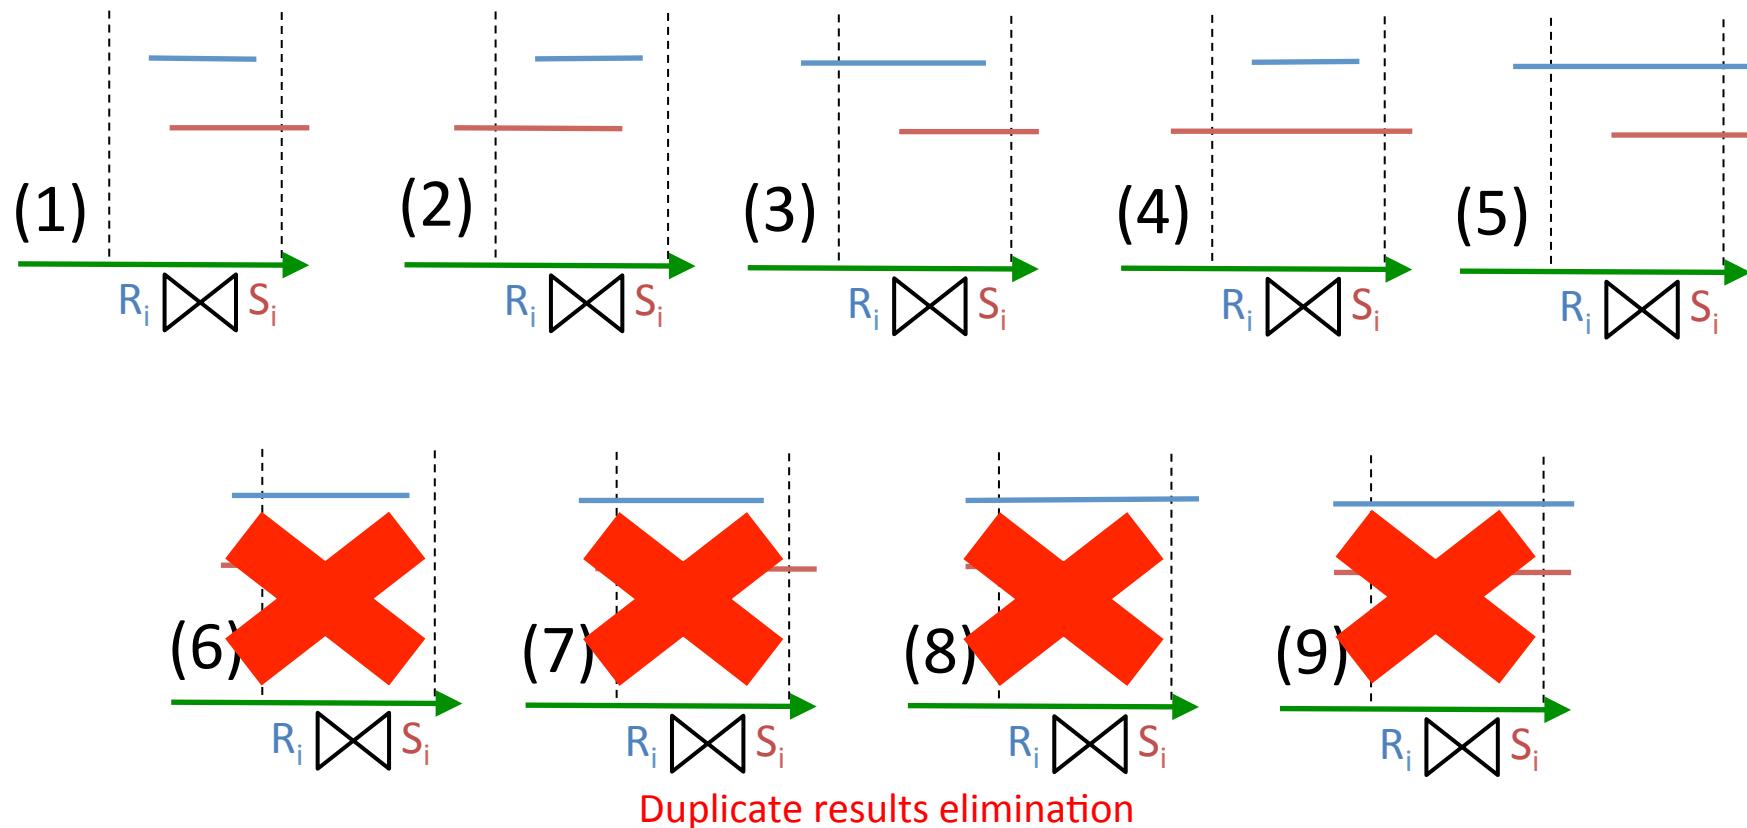
# Mini-joins break down

- 3 types of intervals  $\rightarrow$  9 types of mini-tasks



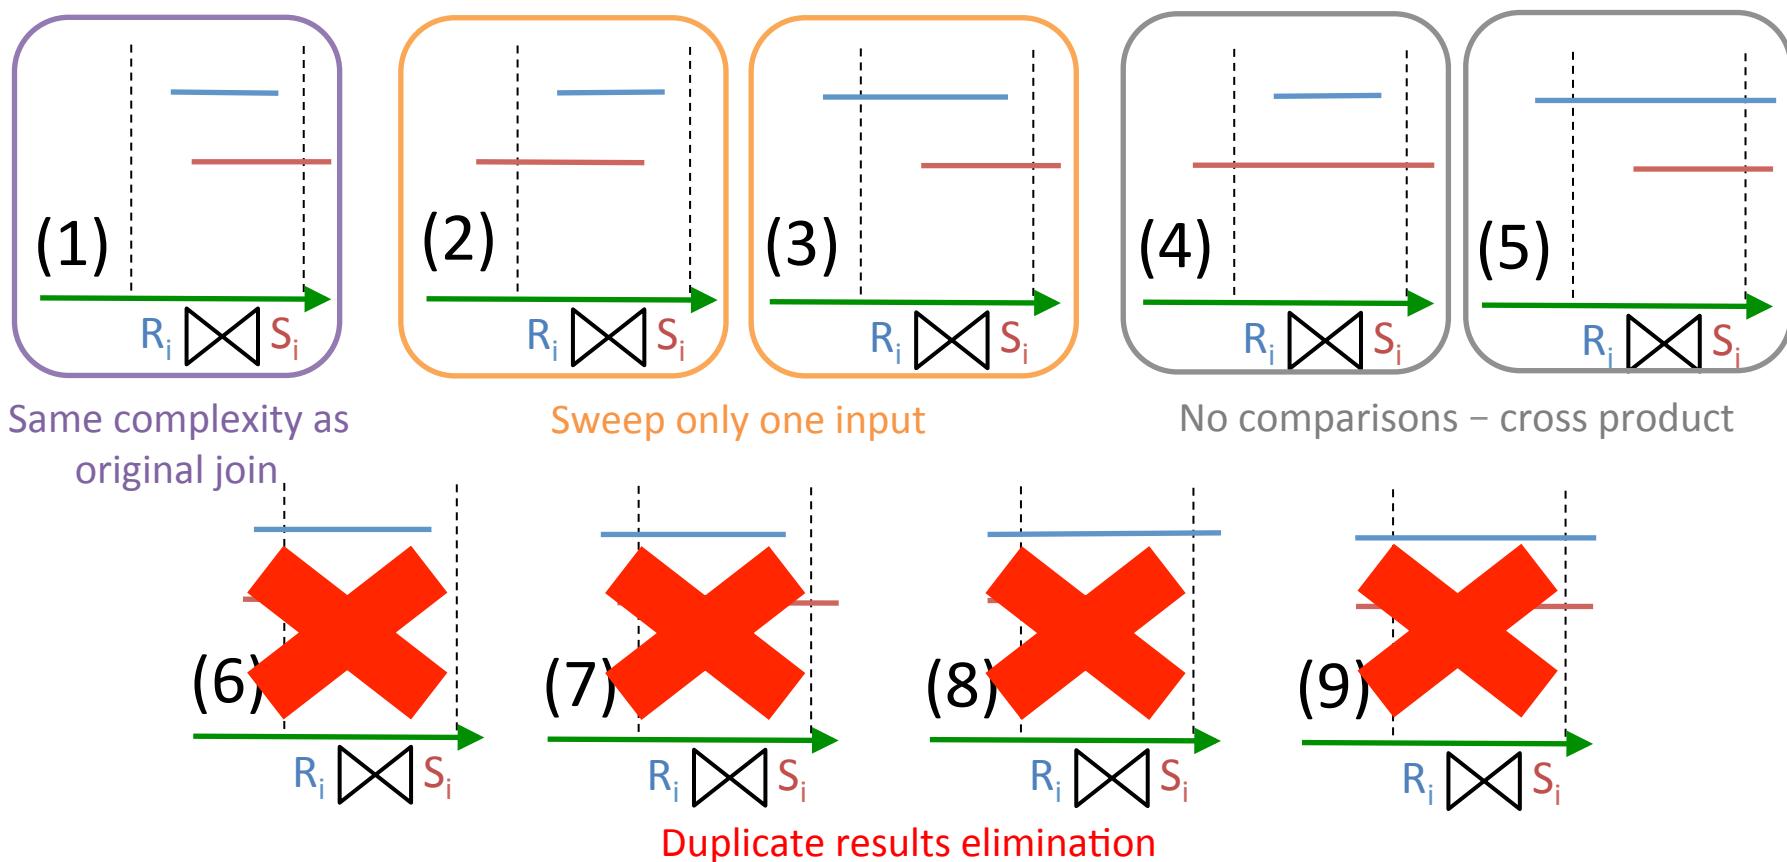
# Mini-joins break down

- 3 types of intervals  $\rightarrow$  9 types of mini-tasks



# Mini-joins break down

- 3 types of intervals  $\rightarrow$  9 types of mini-tasks

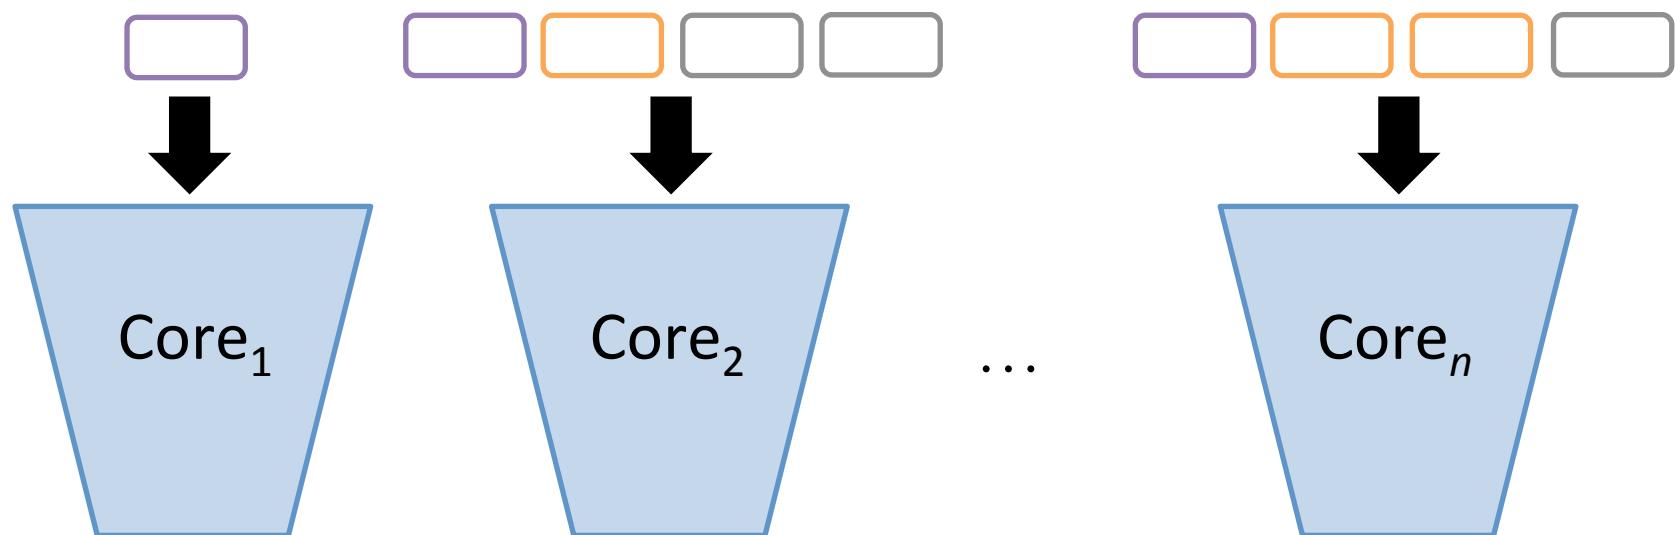


# Greedy scheduling

- Idea
  - Distribute  $1 + 5 \cdot (k-1)$  mini-joins to different cores
  - Evenly distribute load → minimize max load
  - NP-hard problem
    - Greedy approximation algorithm

# Greedy scheduling

- Idea
  - Distribute  $1 + 5 \cdot (k-1)$  mini-joins to different cores
  - Evenly distribute load → minimize max load
  - NP-hard problem
    - Greedy approximation algorithm

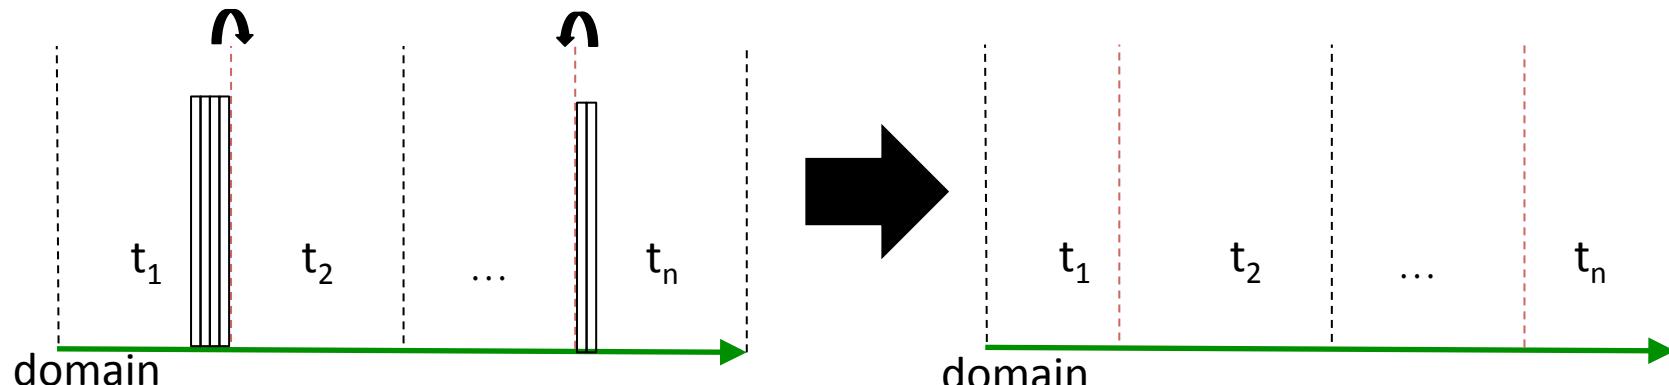


# Adaptive partitioning

- Idea
  - Create an **initial uniform** partitioning
  - Employ a **very fine tiling – granules**
  - **Move load between neighboring tiles**
    - Move granules between **neighboring** tiles
    - Reposition borders of tiles

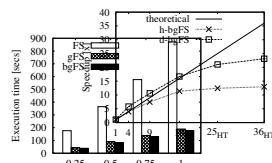
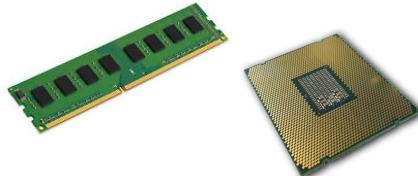
# Adaptive partitioning

- Idea
  - Create an **initial uniform** partitioning
  - Employ a **very fine tiling – granules**
  - **Move load between neighboring tiles**
    - Move granules between **neighboring** tiles
    - Reposition borders of tiles



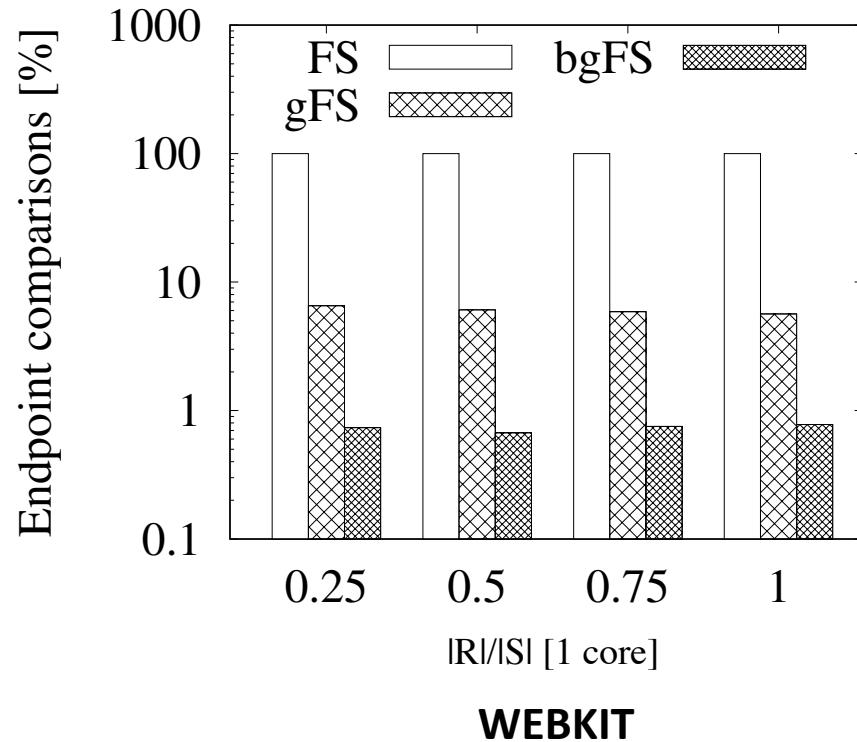
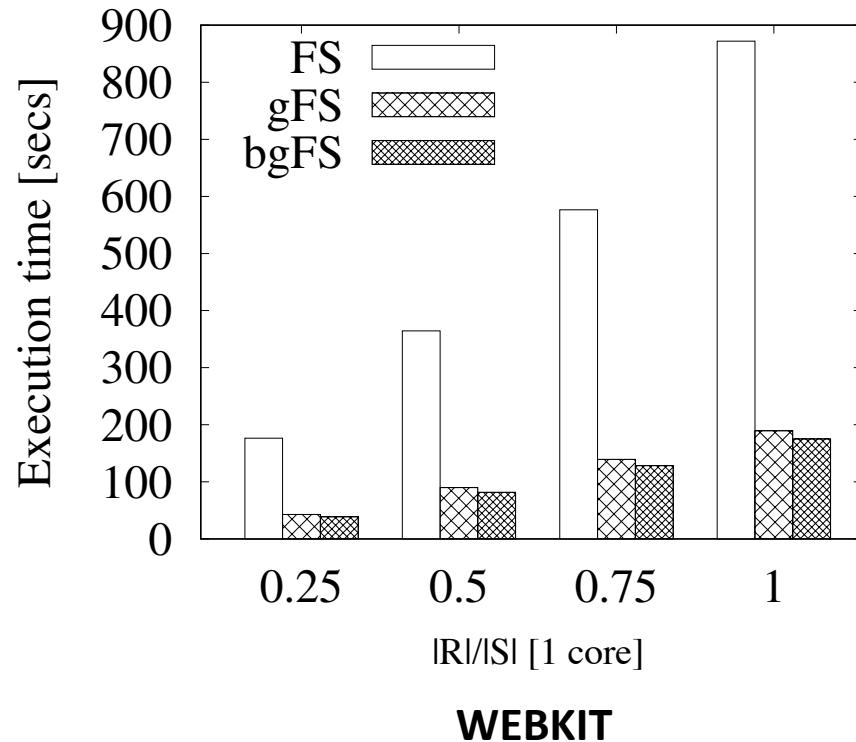
# **EXPERIMENTAL ANALYSIS**

# Setup

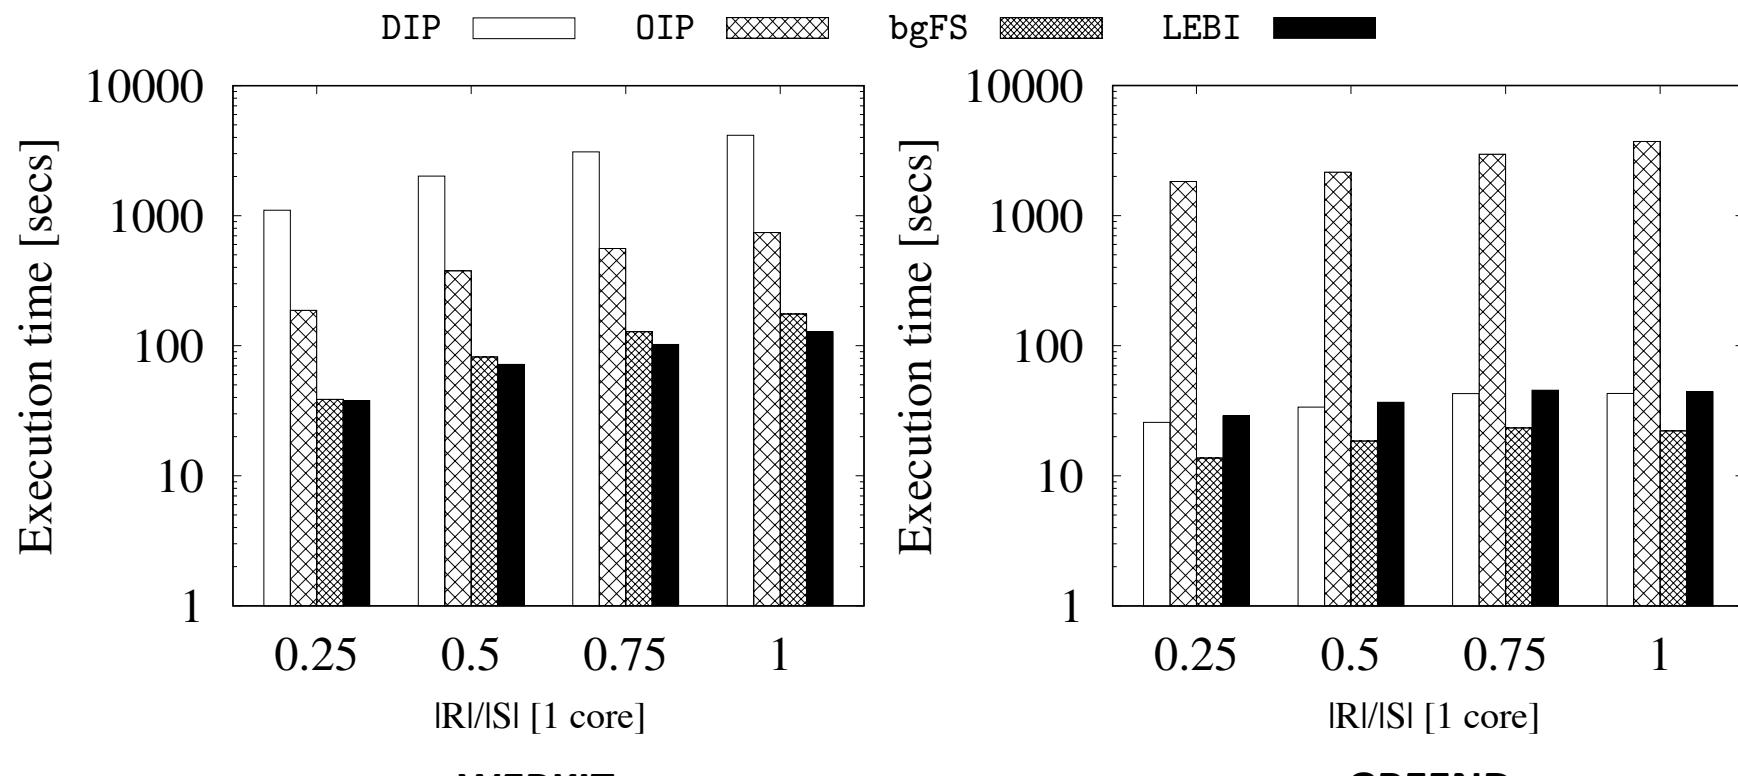


- **Hardware**
  - dual 10-core Intel(R) Xeon(R) CPU E5-2687W v3 @ 3.10 GHz with 128 GBs of RAM
  - Hyper-threading enabled, up to 40 threads
- **Software**
  - Workload [ICDE'16] → XOR of *start*
  - Loop unrolling forced, OpenMP for multi-threading
- **Datasets**
  - WEBKIT git repo, interval = period of time file unchanged
  - BOOKS Aarhus libraries, interval = period of time book lent
  - Synthetic
    - Interval duration follows exponential distribution, uniformly distributed *start* plus peaks
- **Experiments**
  - Execution time, # comparisons, memory footprint
  - Both self joins and non-self joins
  - Vary  $|R|/|S|$ , # cores (threads)

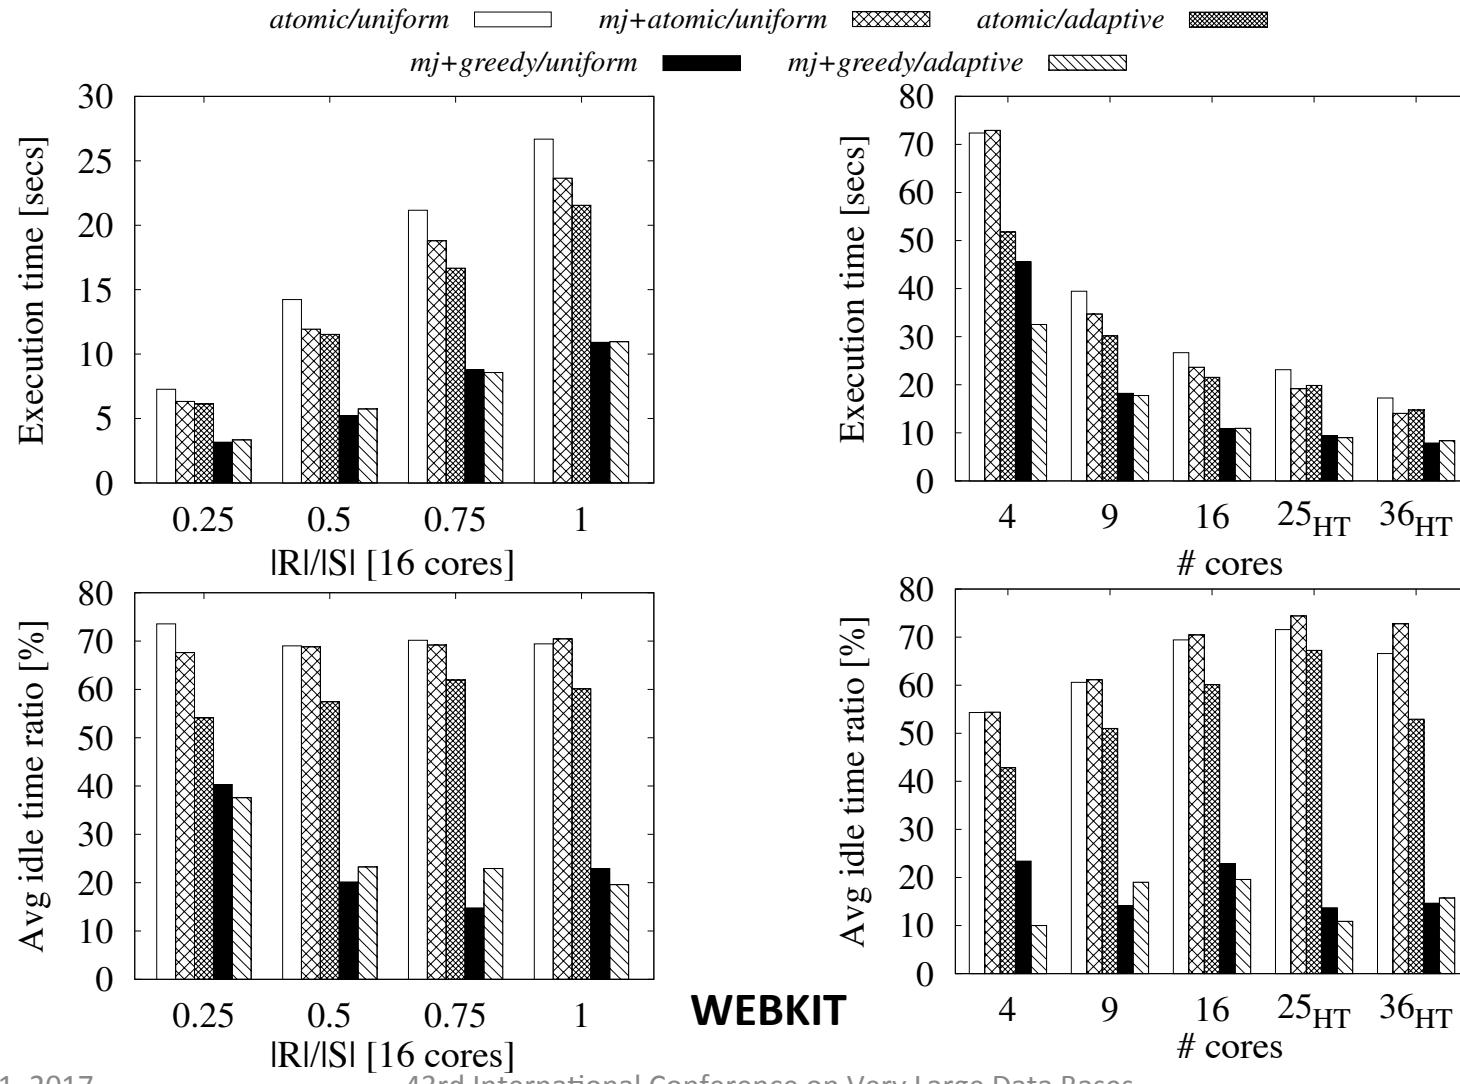
# Optimizing FS



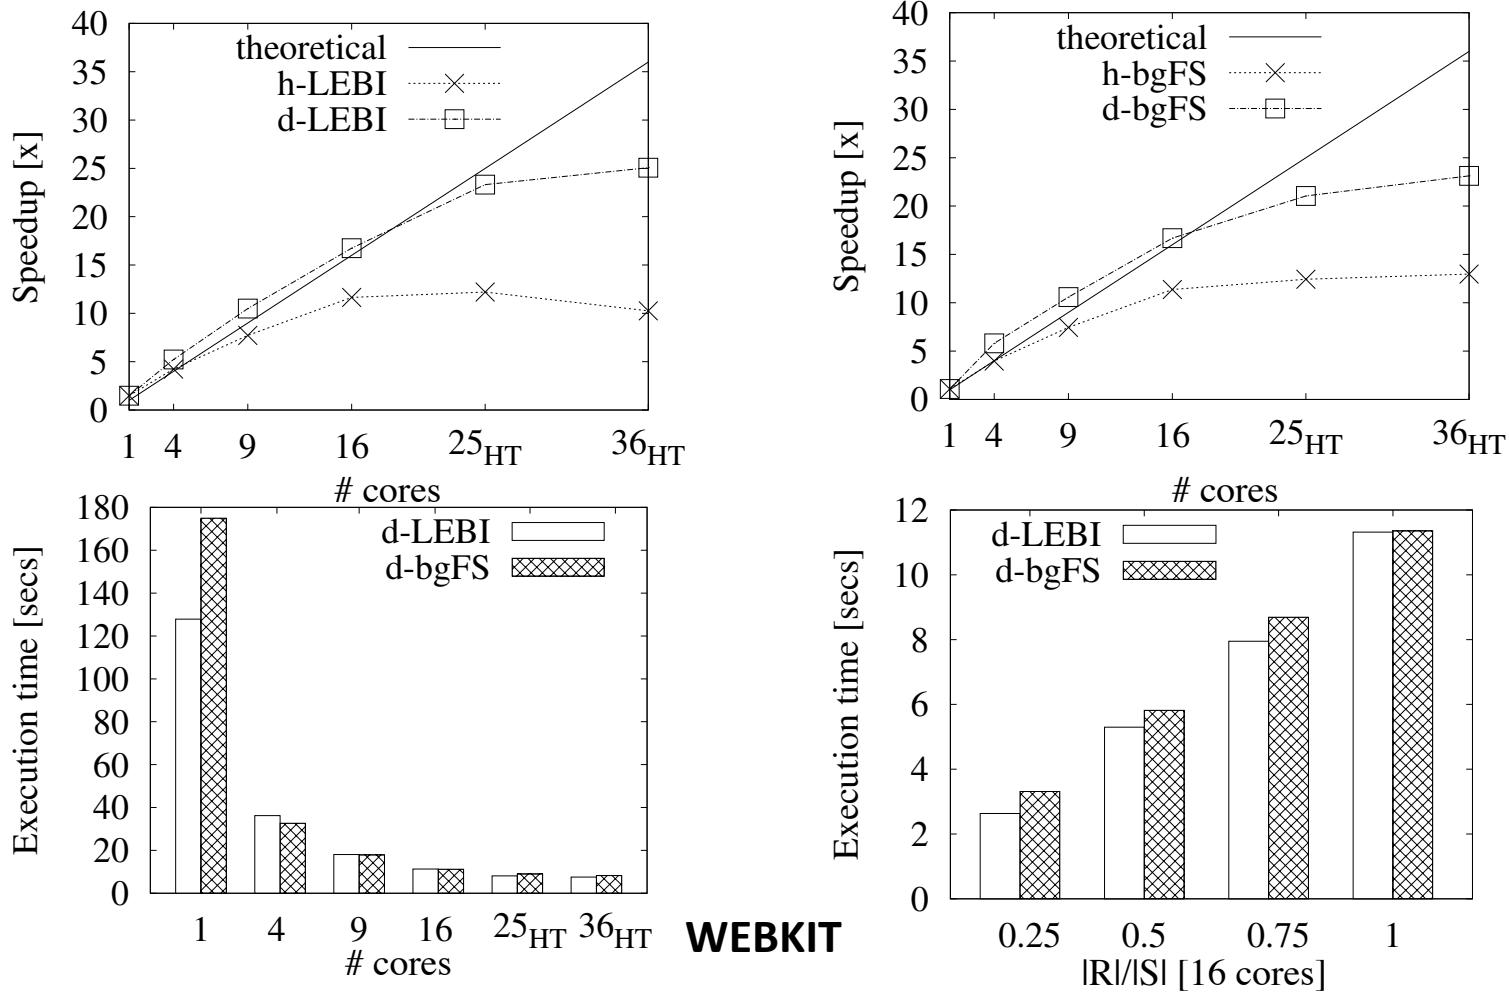
# Single-threaded processing



# Optimizing Domain-based Paradigm



# Parallel processing



# To sum up

- Contributions
  - Efficient evaluation of interval joins
  - Single-threaded processing
    - Optimized bgFS, competitive to state-of-the-art EBI/LEBI
    - Lower memory footprint
  - Parallel processing
    - Novel domain-based partitioning paradigm
    - Higher speedup
- Future work
  - Other types of temporal joins
  - Other types of temporal operators
  - Parallel processing
    - Data-level parallelism, share data between threads

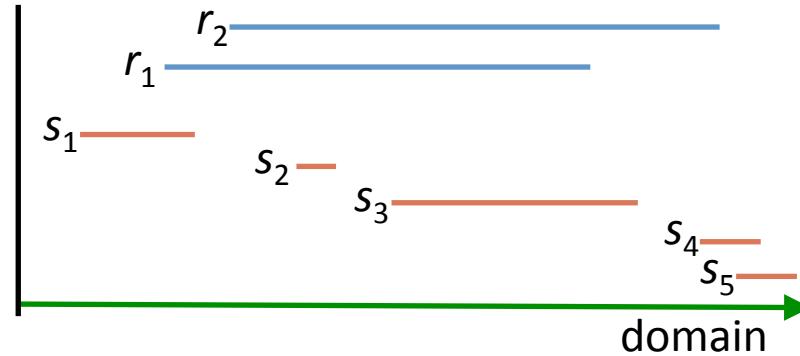
# Questions ?



# **EXTRAS**

# Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]



## Endpoint indices

$$EI^R = \{r_1.start, r_2.start, r_1.end, r_2.end\}$$

$$EI^S = \{s_1.start, s_1.end, s_2.start, s_2.end, s_3.start, s_3.end, s_4.start, s_5.start, s_4.end, s_5.end\}$$

## Active sets

$$A^R = \{\}$$

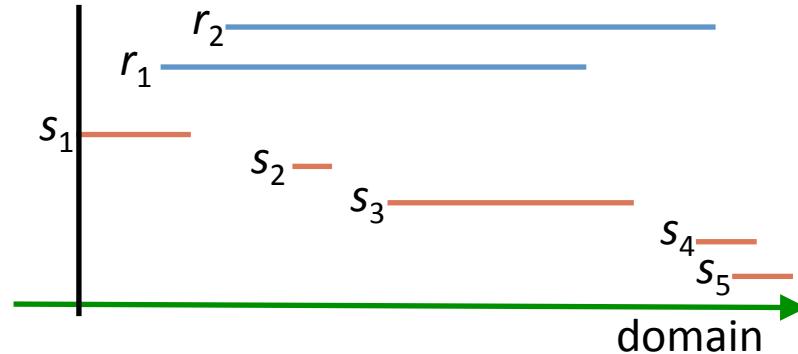
$$A^S = \{\}$$

## Result

$$\{\}$$

# Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]



## Endpoint indices

$$EI^R = \{r_1.start, r_2.start, r_1.end, r_2.end\}$$

$$EI^S = \{\textcolor{red}{s_1.start}, s_1.end, s_2.start, s_2.end, s_3.start, s_3.end, s_4.start, s_5.start, s_4.end, s_5.end\}$$

## Active sets

$$A^R = \{\}$$

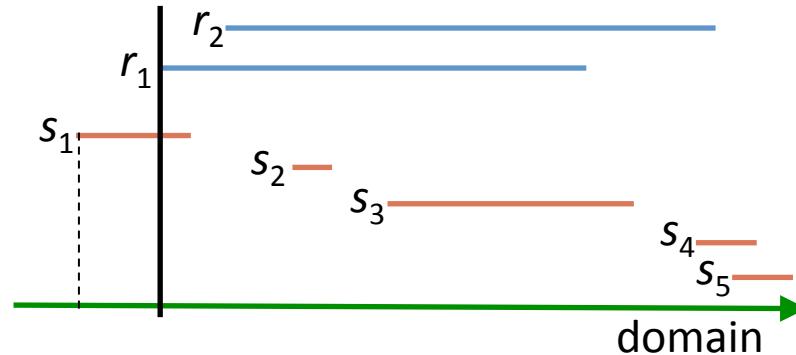
$$A^S = \{\textcolor{red}{s_1}\}$$

## Result

$$\{\}$$

# Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]



## Endpoint indices

$$EI^R = \{r_1.start, r_2.start, r_1.end, r_2.end\}$$

$$EI^S = \{s_1.start, s_1.end, s_2.start, s_2.end, s_3.start, s_3.end, s_4.start, s_5.start, s_4.end, s_5.end\}$$

## Active sets

$$A^R = \{\}$$

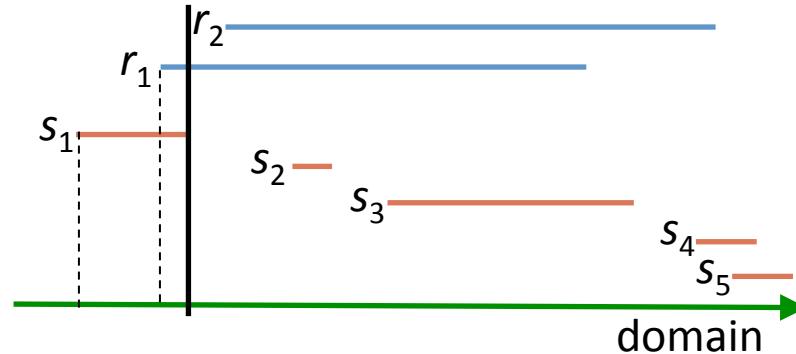
$$A^S = \{s_1\}$$

## Result

$$\{(r_1, s_1)\}$$

# Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]



## Endpoint indices

$$EI^R = \{r_1.start, r_2.start, r_1.end, r_2.end\}$$

$$EI^S = \{s_1.start, s_1.end, s_2.start, s_2.end, s_3.start, s_3.end, s_4.start, s_5.start, s_4.end, s_5.end\}$$

## Active sets

$$A^R = \{\}$$

$$A^S = \cancel{\{s_1\}}$$

## Result

$$\{(r_1, s_1)\}$$

# Endpoint-Based Join (EBI/LEBI)

[Piatov et al., ICDE'16]

## Pros

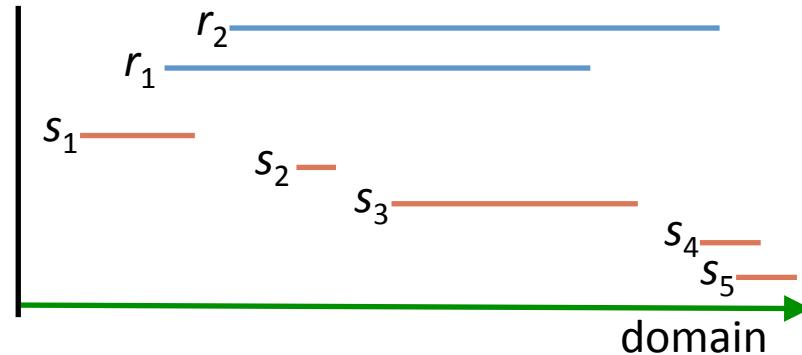
- ✓ No domain-point comparisons when producing results
- ✓ Tailored to modern hardware
- ✓ Main memory cache-aware
- ✓ Fast

## Cons

- ✗ Special data structure needed for active sets, support for efficient updates and scans

# Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]



## Sorted inputs

$$R = \{r_1, r_2\}$$

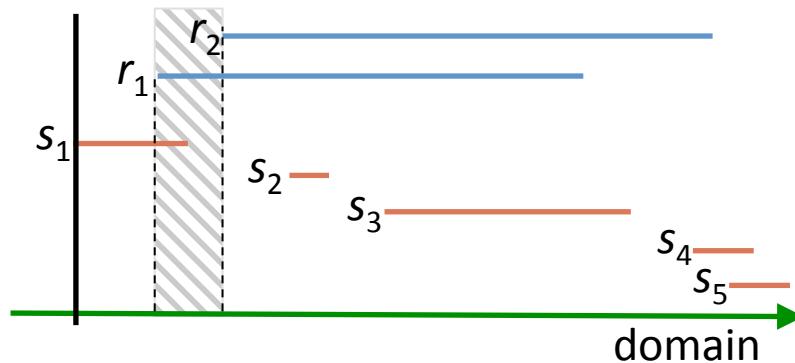
$$S = \{s_1, s_2, s_3, s_4, s_5\}$$

## Result

$$\{\}$$

# Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]



## Sorted inputs

$$R = \{r_1, r_2\}$$

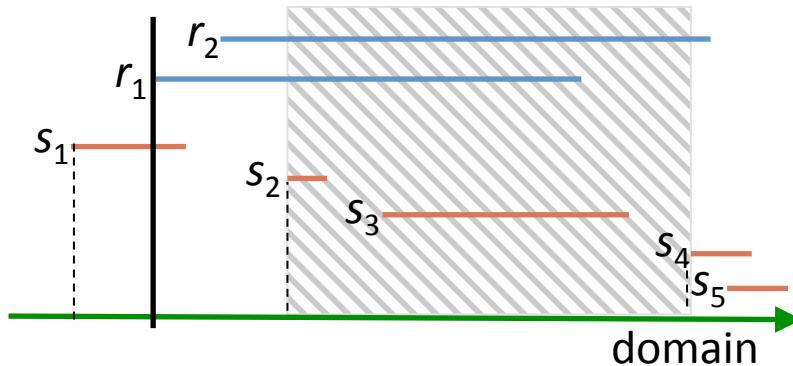
$$S = \{s_1, s_2, s_3, s_4, s_5\}$$

## Result

$$\{(r_1, s_1)\}$$

# Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]



## Sorted inputs

$$R = \{r_1, r_2\}$$

$$S = \{s_1, s_2, s_3, s_4, s_5\}$$

## Result

$$\{(r_1, s_1), (r_1, s_2), (r_1, s_3)\}$$

# Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]

## Pros

- ✓ Simple
- ✓ No special structure needed

## Cons

- ✗ Each join result requires a domain-point comparison,  
 $|R| + |S| + |R \bowtie S|$  comparisons in total

# Forward Scan based (FS)

[Brinkhoff et al., SIGMOD'93]

## Pros

- ✓ Simple
- ✓ No special structure needed

## Cons

- ✗ Each join result requires a domain-point comparison,  
 $|R| + |S| + |R \bowtie S|$  comparisons in total